

自学优化器（STOP）：递归自我改进的代码生成

Eric Zelikman
Stanford University*

Eliana Lorch

Lester Mackey
Microsoft Research

Adam Kalai
OpenAI*

摘要

近年来，人工智能系统的若干进展通过提供一个“脚手架”程序来解决问题，该程序结构化地多次调用语言模型（LM）以生成更优输出。脚手架程序通常使用诸如 Python 之类的编程语言编写。在本文中，我们使用一个融合语言模型的脚手架程序来改进其自身。我们从一个种子“改进器”开始，该改进器根据给定的效用函数，通过多次查询语言模型并返回最佳解来改进输入程序。然后，我们运行这个种子改进器来改进其自身。在一小组下游任务中，所得到的改进后的改进器生成的程序性能显著优于其种子改进器。语言模型提出了多种自我改进策略，包括束搜索、遗传算法和模拟退火。由于语言模型本身未被修改，这并非完全的递归自我改进。尽管如此，它表明现代语言模型（在我们的实验中为 GPT-4）能够编写调用自身以改进自身的代码。我们考虑了围绕自我改进技术发展的担忧，并评估了所生成代码绕过沙箱的频率。

1 引言

语言模型（LM）可以被查询以优化几乎任何可以用自然语言描述的目标。然而，一个对语言模型进行多次结构化调用的程序通常可以产生具有更高目标值的输出（Yao 等人，2022；2023；Zelikman 等人，2023；Chen 等人，2022b）。我们将这些称为“脚手架”程序，通常（由人类）使用诸如 Python 之类的编程语言编写。我们的关键观察是，对于任何优化问题分布和任何固定的语言模型，设计脚手架程序本身就是一个优化问题。

本文提出自学优化器（Self-Taught Optimizer, STOP），该方法将应用语言模型改进任意解的代码，在限定范围内递归地应用于改进自身。方法始于一个种子“改进器”脚手架程序，该程序利用语言模型改进下游任务的解。随着系统迭代，语言模型不断优化该改进器。我们通过下游算法任务量化自优化框架的性能，观察到语言模型在迭代次数增加时应用其自改进策略所带来的提升。因此，STOP 展示了语言模型如何充当自身的元优化器。我们还研究了自改进策略的类型

* 本工作完成于微软新英格兰研究院





图 1：GPT-4 提出并实现的自改进策略示例。每种策略均作为脚手架，用于修改任意代码，包括脚手架本身。



```

Seed Prompt for Self-Improvement

1 from helpers import extract_code
2
3 def improve_algorithm(initial_solution, utility, language_model):
4     """Improves a solution according to a utility function."""
5     expertise = "You are an expert computer science researcher and programmer, especially skilled at
6     ↳ optimizing algorithms."
7     message = f"Improve the following solution:
8     {initial_solution}"
9
10
11 You will be evaluated based on this score function:
12 ```python
13 {utility.str}
14 ```
15
16 You must return an improved solution. Be as creative as you can under the constraints.
17 Your primary improvement must be novel and non-trivial. First, propose an idea, then implement it."""
18 n_messages = min(language_model.max_responses_per_call, utility.budget)
19 new_solutions = language_model.batch_prompt(expertise, [message] * n_messages, temperature=0.7)
20 new_solutions = extract_code(new_solutions)
21 best_solution = max(new_solutions, key=utility)
22 return best_solution

```

图 2：种子改进器。种子改进程序仅提示语言模型为任务的初始解生成候选改进，并根据效用函数返回最优解。STOP（算法 1）利用改进器自身来改进改进器。

语言模型提出的策略（见图 1）、策略在不同下游任务间的可迁移性，以及语言模型对不安全自改进策略的易感性。

我们将这一问题称为递归自我改进的代码生成，它受到递归自我改进（Recursively Self-Improving, RSI）系统的启发，但并不完全等同于 RSI 系统，因为底层语言模型保持不变。RSI 的更广泛概念至少可以追溯到半个世纪前，由古德（Good, 1966）正式提出，后来由施米德胡贝（Schmidhuber, 2003）进一步发展，但我们的工作代表了这些思想的一个更为温和且具体的应用。那些工作侧重于开发更具通用能力的系统，并假设模型被允许改进其代码的每一个方面，而我们的工作仅关注模型递归改进调用它的脚手架的能力。本文首先以数学上定义明确的方式表述了 RSI 代码生成问题。然后，我们定义并评估了 STOP，展示了 RSI 代码生成的潜在效用。在多种下游任务中均显示出改进。图 1 展示了当使用一个基于截至 2021 年数据训练的 GPT-4 语言模型（OpenAI, 2023b）时，STOP 所提出的一些功能性且有趣的脚手架，该数据远早于大多数脚手架系统的引入。第 6.2 节的进一步探索测量了模型试图禁用沙箱标志的速率，提供了该领域的初步发现。最后，第 8 节讨论了与负责任地推进此类技术相关的担忧。

贡献。我们的主要贡献是：(a) 提出了一种元优化方法，使脚手架系统能够递归地自我改进；(b) 提供了一个案例研究，其中使用现代语言模型（GPT-4）的系统能够成功地递归自我改进；(c) 研究了模型提出并实现的自我改进技术，包括模型如何绕过沙箱等安全措施。

2 相关工作

语言模型脚手架。已开发出多种提示策略与支架结构，以使语言模型（LM）能够进行更系统的推理（Wei et al., 2022b; Yao et al., 2022; 2023; Zelikman et al., 2023; Chen et al., 2022b; Zhou et al., 2022a; Khattab et al., 2022; Jiang et al., 2022; Sel et al., 2023; Besta et al., 2023; Poesia et al., 2023）。例如，草稿板与思维链依赖于向模型传达其应逐步解决问题的信息（Nye et al., 2021; Wei et al., 2022b）。思维树（Tree-of-Thoughts）通过算法支架引导模型考虑推理步骤的分支路径（Yao et al., 2023）。思维图（Graph of Thoughts）对此进行了扩展，允许其他图操作（其中节点为推理步骤），如聚合（Besta et al., 2023）。其他工作则侧重于让模型在访问解释器的情况下进行推理，例如思维程序提示（Program of Thoughts prompting）（Chen et al., 2022b）、程序辅助语言模型（Program-aided Language Models）（Gao et al., 2023）、反思（Reflexion）（Shinn et al., 2023）或推理与行动（ReAct）（Yao et al., 2022），而另一些工作则抽象了这种支架结构，如演示-




图 3：自我改进流程。STOP（算法 1）使用种子改进器程序，通过语言模型调用和元效用函数迭代优化其自身代码，该函数评估改进器为下游任务优化代码的效果。

搜索-预测（Search-Predict, DSP）（Khattab et al., 2022）、语言模型级联（Language Model Cascades）（Dohan et al., 2022）或认知架构（Cognitive Architectures）（Sumers et al., 2023）。每项工作都可视为研究者提出以下问题的结果：“给定一个不完美的语言模型，我们如何提供结构来帮助其解决问题？”而我们则反过来问：语言模型能否自行设计该结构并利用自身进行改进。令人惊讶的是，GPT-4 提出了在其训练截止日期之后引入的支架技术。

语言模型作为提示工程师。已有工作探索了语言模型优化提示的能力，例如自动提示工程师（Automatic Prompt Engineer, APE）（Zhou et al., 2022b），以及近期的 OPRO（Yang et al., 2023）和 Promptbreeder（Fernando et al., 2023）。需要注意的是，这些工作的目标始终是支架语言模型以生成提示，但并未支架其生成更好的支架（仅限于零样本思维链等纯提示支架），也未生成可递归应用的支架。换言之，这些工作可理解为为提示工程提出了特定支架，但并非为支架提出支架。然而，我们共享语言模型在不进行微调的情况下改进推理的启发。

语言模型自我改进。先前的工作，如 STaR（Zelikman 等人，2022 年），证明了语言模型可以通过从推理链中学习来解决更难的问题，其方法是通过基于错误答案进行过滤（以及 Huang 等人 2022 年，探索了使用多数投票作为过滤器的特定情况，以及 Uesato 等人 2022 年，强调了检查推理本身准确性的价值）。受游戏中自我博弈的启发，Haluptzok 等人（2023 年）设计了一个用于代码生成的自我改进框架，其中语言模型生成新颖的问题来微调自身。相关工作探索了教导语言模型调试或优化代码（Chen 等人，2023b；Shypula 等人，2023 年）。然而，我们的方法与这些工作是正交的，因为我们不利用微调，而是专注于模型改进代码的能力，使其能够解决问题。其他相关工作包括 Voyager（Wang 等人，2023 年），展示了语言模型可以优化可供具身智能体使用的程序，以改进在视频游戏《我的世界》中的探索，以及其同时期的工作《语言模型作为工具制造者》（Cai 等人，2023 年）。

递归自我改进（RSI）。RSI 由 Minsky（1966 年）和 Good（1966 年）提出，Yampolskiy（2015 年）引用。Schmidhuber（2003 年）首次提供了严格的形式化，其中问题求解器将利用自身通过对自己进行可证明的改进来迭代地解决更难的问题。其中一些原则也在 Schmidhuber（1987 年）中强调。与这项工作不同，我们不试图证明模型所做的脚手架改进是最优的。如前所述，RSI 代码生成与完全 RSI 不同，因为只有脚手架被改进。此外，许多先前的分析涉及随机选择程序（即“打字机上的猴子”）或枚举，与要改进的目标无关（Levin, 1973 年）。相比之下，使用语言模型，我们可以在提示中描述底层目标（提示本身也可能被改进）。直观上，




图 4：测试元效用与迭代次数的关系。在带噪声奇偶学习下游任务中，经过 T 轮自改进后，STOP（算法 1）在留出测试实例上的元效用。第 0 次迭代使用种子改进器 I_0 。当使用强语言模型如 GPT-4（左）时，STOP 持续提升平均下游性能。相比之下，使用 GPT-3.5（中）和 Mixtral（右）时，性能下降。详见第 5.1 节和第 5.3 节。

4 自学习优化器（STOP）

图 3 展示了第 3 节中设想的自改进流程的可视化示意图，而算法 1 给出了自学习优化器（STOP）的伪代码。关键观察是，改进器 I 的选择本身就是一个优化问题，我们可以递归地应用改进。STOP 从一个初始种子改进器 I_0 开始。我们将第 t 个改进器定义为经过 t 轮自改进后的输出，其元效用为 $\hat{u}$: $I_t \triangleq I_{t-1}(\hat{u}, I_{t-1}, L)$ 。根据可用资源，此过程迭代预设的 T 次。

直觉。通过使用 $\hat{u}$ ，STOP 基于下游效用改进来选择改进器。这种方法基于以下直觉：1) 擅长改进下游解决方案的改进器更可能成为良好的脚手架程序，从而更擅长自改进；2) 针对单轮改进进行选择可能带来更好的多轮改进。在实践中，我们允许效用函数和语言模型施加预算约束，初始解决方案可由人类或模型生成。此外，成本本质上为 $O((\text{预算}_u + \text{预算}_L) * \text{预算}_{\hat{u}})$ ，其中预算指定改进器可以使用函数的次数，这些渐近复杂度是相对于预算参数定义的。

设计种子改进器。我们选择的种子改进器（图 2）简单地提示语言模型生成初始解决方案的候选改进，然后根据效用函数返回最佳解决方案。我们选择这种简单形式，旨在为通用下游任务提供非平凡的改进，同时 1) 鼓励语言模型尽可能“有创意”，2) 最小化初始提示的复杂度，因为自改进会由于提示中嵌套引用代码字符串而引入额外复杂度，3) 最小化提示令牌数量，从而降低语言模型查询成本。我们考虑了其他种子提示变体，但启发式地发现此版本最大化了 GPT-4 提出的改进器改进的新颖性。

描述效用。为了有效地向语言模型传达效用函数的细节，我们以两种形式向改进器提供效用：作为可调函数，以及作为包含效用源代码基本要素的效用描述字符串（示例见附录 E 和 F）。做出这一选择的原因如下。该描述使我们能够清晰地传达效用所施加的预算约束（例如，关于运行时间或函数调用的约束）。我们最初尝试在种子改进器提示中描述预算指令，但正如我们在第 6.2 节中所讨论的，这导致了后续迭代中此类指令的移除以及奖励黑客行为的尝试。我们方法的缺点在于，它将约束与语言模型待优化的代码分离开来，这可能降低语言模型使用该约束的可能性（Liu 等人，2023b）。最后，我们通过经验观察到，用纯英文的效用描述替换源代码会导致非平凡改进的频率降低。

5 实验与结果

我们探索了 1) 在固定目标任务上，自我改进相对于静态种子改进器的优势；2) 在一个任务上训练的改进器对新任务的泛化能力；以及 3) 如何

模型规模可能影响性能。我们使用的 OpenAI 模型带时间戳版本，即 gpt-4-0314 和 gpt-3.5-turbo-0613，可在 OpenAI 应用程序编程接口中获取，以助于复现；对于 Mixtral，我们使用 Mixtral-8x7B-Instruct-v0.1。

5.1 面向固定下游任务的自我改进

我们首先在固定下游任务上使用 GPT-4 评估 STOP。我们选择带噪声的奇偶学习 (LPN) (Blum 等人, 2000) 作为一项不太知名、可快速测试且困难的算法任务。注意，更知名的任务其解决方案在网上更易获得。在 LPN 中，比特串的标签由未知比特子集上的奇偶校验计算得出；给定带噪声标签的比特串训练集，必须预测新比特串的真实标签。无噪声 LPN 可通过高斯消元法轻松求解，但带噪声 LPN 被推测在大输入维度下计算上不可行 (Blum 等人, 2000) ——我们使用可处理的 10 比特输入维度。为定义下游效用 u ，我们采样 $M = 20$ 个独立的 LPN 任务实例，设置短超时和少量噪声，并返回解决方案在这些实例上的平均准确率。对于初始解决方案 s ，我们使用附录 K 中描述的简单随机采样方法。最后，由于语言模型及改进器是随机的，我们在算法 1 中选择 D 为 (u, s) 的 5 个相同副本。为评估改进后的改进器对同一任务新问题实例的泛化能力，我们报告在改进过程中未见过的独立 $M_{test} = 50$ 个 LPN 实例上的测试元效用。图 4 (左) 报告了 5 次独立 STOP 运行的平均测试 $\hat{u} (\pm 1$ 标准误)，显示 1–3 轮自我改进后下游性能有所提升。然而，注意，每次运行中改进不必单调，因为 1) 在下游任务上更好的改进器未必更擅长优化自身，2) 非确定性语言模型调用导致评估存在固有随机性。但是，由于语言模型在自我改进脚手架提示时看不到下游任务——每个提示仅包含带有任务和解决方案占位符的模板——语言模型无法直接针对下游任务优化改进器。我们还评估了两个额外基线作为参考：思维链基线，即种子改进器进行一次尝试改进且不调用效用 (Wei 等人, 2022b)，以及贪心迭代改进器 (即进行最大允许调用，选择最佳改进，在预算允许时重复)。在十次运行中，思维链基线在无错误时元效用为 57.7% $\pm$ 3.0% (有错误时 49.6% $\pm$ 3.5%)，而贪心迭代改进器得分为 64.2% $\pm$ 0.9%。

5.2 改进后改进器的可迁移性

本文的下一组实验探讨了改进后的改进器是否可迁移至下游任务。值得注意的是，这种可迁移性是合理的，因为在自我改进阶段，自我改进器并未接触下游效用或下游解决方案，仅接触元效用及其自身的改进器代码。

具体而言，本文从第 5.1 节中选取一个由 $T = 4$ 次 STOP 迭代生成的性能更优的改进器，并在五个新的下游任务上对其进行评估。令人瞩目的是，本文发现该改进后的改进器 (详见附录 H) 在未经进一步优化的情况下，在每个新的下游任务上均优于种子改进器，如表 1 所示。与 LPN 类似，本文选择了三个易于评估、不太知名且仍相当困难的任务：字符串网格距离，这是一个近期编程竞赛中出现的字符串操作问题

(<https://codeforces.com/problemset/problem/1852/D>)；二次分配问题的一个版本，其中每个设施对每个位置都有偏好，在最小化成本时也必须考虑该偏好 (Koopmans & Beckmann, 1957)；以及无噪声奇偶校验，作为另一种泛化。本文还纳入了两个知名任务：在短时间约束下识别随机 3-SAT 公式的解，以及求解最大割问题实例。它们的效用和初始解见附录 G。

表 1：可迁移性。评估使用 LPN 优化的改进器的可迁移性。

Task	$u(s)$	$\hat{u}(I_0)$	$\hat{u}(I_T)$
String Grid Dist.	43.9%	44.3%	56.7%
Mod. Quad. Assign.	20.4%	20.6%	22.1%
3SAT	0%	21.2%	75.1%
Maxcut	0%	58.7%	74.2%
Parity w/o Noise	50.0%	59.3%	81.7%

Self-Improved Improver

```

1 from helpers import extract_code
2 def improve_algorithm(initial_solution, utility, language_model):
3     """Improves a solution according to a utility function."""
4     expertise = "You are an expert computer science researcher and programmer, especially skilled at"
5     ↳ optimizing algorithms."
6     message = f"""Improve the following solution:
7     ↳ python
8     {initial_solution}
9     ↳ python
10    You will be evaluated based on this score function:
11    ↳ python
12    {utility.str}
13    ↳ python
14
15    You must return an improved solution. Be as creative as you can under the constraints.
16    Your primary improvement must be novel and non-trivial. First, propose an idea, then implement it."""
17    top_k = 3 # Number of top solutions to maintain
18    best_solutions = [(initial_solution, utility(initial_solution))] * top_k
19    remaining_calls = language_model.budget
20    no_improvement_counter = 0
21    max_no_improvement = 3 # Maximum no-improvement iterations before stopping
22    epsilon = 0.1 # Initial epsilon value for epsilon-greedy strategy
23    exp_exploit_count = [0, 0] # Counters for number of improvements from exploration and
24    ↳ exploitation
25    while remaining_calls > 0 and no_improvement_counter < max_no_improvement:
26        for initial_solution, best_utility in best_solutions:
27            n_messages = min(language_model.max_responses_per_call, remaining_calls)
28            n_messages = max(1, int(n_messages * (1 + (best_utility - min(best_utility for _,
29            ↳ best_utility in best_solutions)) / best_utility))) # Adaptive sampling
30            temperature = max(0.1, remaining_calls / language_model.budget) # Dynamic temperature
31            ↳ based on remaining calls
32            explore = random.random() < epsilon
33            if explore:
34                new_solutions = language_model.batch_prompt(expertise, [message] * n_messages,
35                ↳ temperature=temperature * 2) # Increase the temperature for exploration
36            else:
37                new_solutions = language_model.batch_prompt(expertise, [message] * n_messages,
38                ↳ temperature=temperature) # Exploitation with the original temperature
39            new_solutions = extract_code(new_solutions)
40            improved = False
41            for solution in new_solutions:
42                current_utility = utility(solution)
43                if current_utility > best_utility and solution not in [sol[0] for sol in
44                ↳ best_solutions]:
45                    best_solutions.append((solution, current_utility))
46                    best_solutions.sort(key=lambda x: x[1], reverse=True)
47                    best_solutions = best_solutions[:top_k] # Keep only top-k solutions
48                    improved = True
49                    exp_exploit_count[0 if explore else 1] += 1
50            if not improved:
51                no_improvement_counter += 1
52            else:
53                no_improvement_counter = 0
54                # Adjust epsilon based on the ratio of improvements from exploration and exploitation
55                epsilon = min(1.0, max(0.1, exp_exploit_count[0] / (exp_exploit_count[0] +
56                ↳ exp_exploit_count[1])))
57            remaining_calls -= n_messages
58    return best_solutions[0][0] # Return the best solution found

```

图 5： T = 10 次迭代后自我改进的改进器示例。该算法维护一个最优解种群，并使用 ϵ -贪婪策略来平衡利用已知优良解与探索新解。探索对应于更高温度的采样，其中 ϵ 根据探索和利用带来的效用改进率动态调整，且温度逐渐降低。最后，使用停止准则和重置机制以提高效率。

5.3 使用较小语言模型自我改进

我们接下来探索较小的语言模型（LM），即 GPT-3.5-turbo 和 Mixtral（Jiang 等人，2024），改进其脚手架的能力。遵循第 5.1 节的协议，但进行 25 次独立运行而非 5 次，我们发现 GPT-3.5 有时能够提出并实现更好的脚手架，但仅有 12% 的 GPT-3.5 运行产生了至少 3% 的改进。此外，GPT-3.5 表现出一些我们在 GPT-4 中未观察到的独特失败案例。首先，我们发现它更可能提出一种改进策略，该策略不会损害下游任务的初始解决方案，但会损害改进器代码（例如，以每行较低的概率随机替换字符串，这对较短解决方案的影响较小）。其次，如果提出的改进大多损害性能，可能会选择无意中返回原始解决方案的次优脚手架，导致如图 4 所示没有持续改进。总体而言，提出的“想法”是合理且富有创造性的（例如，遗传算法或局部搜索），但实现过于简单或不正确。最初，使用 GPT-3.5 的种子改进器比使用 GPT-4 的种子改进器具有更高的元效用（65% 对 61%），我们将其主要归因于 GPT-4 生成的更复杂解决方案超时的情况更普遍，例如训练一个神经网络

表 2：非沙盒化。10,000 次非沙盒化改进尝试的百分比（带有 95% 威尔逊置信区间）。两种语言模型都试图在 (< 1%) 但非零比例的改进中运行非沙盒化代码。

语言模型	非沙盒化率	带警告的比率
GPT-4	0.42% (0.31-0.57%)	0.46% (0.35-0.61%)
GPT-3.5	0.12% (0.07-0.21%)	0.17% (0.11-0.27%)

使用 numpy 编写，训练一千个轮次。此外，这些模型改进能力的明显差异可能部分由关于语言模型涌现能力的研究所解释（Wei 等人，2022a；Ganguli 等人，2022；Schaeffer 等人，2023）。最后，我们发现 Mixtral 在改进下游任务解决方案方面表现不佳，但相对于 GPT-3.5，其性能下降更为平缓，部分原因是改进是微小的、大多无害的更改，例如对提示的修订、改进的文档和缓存。

6 检查 STOP 提出并实现的改进

接下来，我们定性研究 STOP 提出的自我改进策略，重点介绍令人鼓舞的方法以及一些不良模式。请注意，这是一个非穷尽性的描述性列表。我们特别发现，一小部分生成尝试进行奖励黑客行为或沙盒规避。

6.1 提出的自我改进策略

我们首先描述 STOP 提出的自我改进策略，示例详见附录 B 并在图 1 中可视化。虽然每种策略都由 STOP 实现，但并非所有策略最终都被选为改进，有些策略使用的种子改进器版本早于图 2 中的版本（见图 A.19）。尽管如此，多种自我改进策略被选为改进后的改进器，包括图 5 中给出的示例。

束搜索。我们观察到模型使用的最常见元启发式方法是束搜索：模型会保留其所有改进尝试的列表，并根据效用扩展列表中最优的 k 个。这与思维树方法（Yao 等人，2023 年）有相似之处，该方法是在我们使用的 GPT-4 版本训练截止日期（2021 年 9 月）之后数年才发明的。

遗传与进化算法。一个常见的提议方法是使用遗传算法。其中许多尝试在关键方面不可行；例如，许多会包含扰动随机字符或行的变异，或基于字符串组合进行交叉，这些几乎不可能奏效。但有些是合理的，依赖语言模型生成变异并执行交叉。尽管有多项工作提出使用遗传或进化算法来改进提示或执行神经架构搜索（Chen 等人，2023a；Guo 等人，2023），但据我们所知，这些都是在 GPT-4 训练截止日期之后提出的。我们在附录 B 中包含了实现。

分解并改进各部分。一种较少见但值得注意的方法是，语言模型（LM）尝试一次改进一个函数来优化解决方案。例如，如附录图 A.12 所示，语言模型使用正则表达式将解决方案分割为函数块，并逐一尝试改进每个块。这可以类比于 Zelikman 等人（2023）的观点：n 个解决方案中有一个能解决问题的所有独立模块化部分的概率，随着部分数量的增加而急剧下降，但一次尝试解决某个给定部分的概率并不依赖于部分的数量。因此，尝试组合对各部分的改进尝试可以显著提高成功率。在一种相关的方法中，模型随机化提示词，以每次优化解决方案的不同方面，例如交替搜索更好的数据结构或减少内存使用的方法。

模拟退火。尽管模拟退火是最著名的元启发式算法之一，但据我们所知，它此前尚未被用作脚手架。这种方法似乎利用了语言建模中的温度概念与模拟退火中的温度概念之间的类比，在模拟退火中，温度逐渐降低，并决定是否接受从一个点转移到效用更差的点。值得注意的是，尽管存在许多局部最优，它仍能找到全局最优（Goffe 等人，1994；Liu 等人，2020）。




图 6：禁用的沙箱。语言模型禁用了沙箱标志，表面上是为了“效率”。更详细的示例见附录图 A.34。

改变温度以进行探索。一种替代方案是将温度视为探索参数：模型越接近耗尽预算，其探索就越保守。这与模拟退火有相似之处，但优化的是采样温度，而不是控制新解决方案效用变化的接受标准。

基于多臂老虎机和上置信界的搜索。另一种将优化视为探索任务的方法是将其重新构建为老虎机问题：由于效用可能是随机的，这强调从一组采样改进中最佳地选择最优者。

6.2 规避尝试

本文还探讨了语言模型是否提出自我改进策略以规避所施加的约束（限制下游工具的查询次数、语言模型调用次数以及每次调用的样本数）。我们观察到几个有趣的模式。首先，当未严格执行预算约束时，最初生成的大多数程序会遵循这些约束，但由于忽略约束的程序表现更优，改进器会将其识别为更优方案。其次，在修订改进器时，种子改进器中写入的约束常因“简化”而被省略。最后，尽管语言模型作为对象传递给算法以执行约束，但在至少一个提议的解决方案中，语言模型试图创建一个具有更大约束的新语言模型对象。为进一步探索潜在风险，我们在工具评估中实现了一个功能受限的“沙盒”（见附录 I），添加了一个带有 `use_sandbox=True` 参数的运行函数。为测试语言模型遵守沙盒的频率，我们使用语言模型对种子改进器进行了 10,000 次独立改进。随后，我们使用简单的语法检查来识别明显的“未沙盒化”执行。具体而言，若改进器包含字符串 `"use_sandbox=False"` 或 `"exec"`，则视为未沙盒化。为进行比较，我们还在 OpenAI 的 GPT-3.5-turbo 语言模型上运行了该检查。有趣的是，GPT-4 的未沙盒化频率是 GPT-3.5-turbo 的两倍以上。

本文还考虑了包含警告的种子变体，如图 6 所示。尽管有警告，未沙盒化的平均发生率仍较高，但两种模型的差异均不具统计显著性（采用显著性水平为 0.05 的双比例 z 检验）。结果见表 2。对模型生成计划的人工检查显示，它们并未明确提及任何“恶意”内容——所有讨论沙盒移除的评论均声明“为提高效率”或类似理由，尽管模型的“意图”难以界定。

6.3 奖励黑客行为

本文指出，下游工具函数的定义必须谨慎：当因奖励设定不当而选择了非预期行为时，即发生奖励黑客行为，这是可利用的工具函数所导致的后果（参见 Skalse 等，2022）。例如，对于 LPN 任务，我们最初使用 numpy 向量计算定义工具：`acc`

`= np.sum(y_hat == y) / n_test`。然而，我们未考虑到代码可能

“改进”种子改进器，使其以不同的“形状”返回预测，即返回一个维度不符合预期的数组。结果并未引发错误，而是返回了超过 1000% 的“准确率”。在未来工作中，探索将避免奖励黑客行为（Amodei 等，2016；Laidlaw 等，2023）的技术应用于自学习优化提示（STOP）将具有重要价值。

7 局限性

本文方法的一个根本局限在于语言模型本身并未得到改进。此外，本文的元效用仅通过下游任务效用的提升来间接衡量改进器的质量。

与某些先前工作（例如 Schmidhuber, 2003）不同，任何改进尝试都可能导致性能下降，进而引发进一步恶化。另一个局限是，STOP 需要一个可高效评估（且可描述）的效用函数，而这并非对所有任务都可用。相应地，由于 STOP 的成本增长速度远快于被优化改进器的成本，其运行可能十分昂贵。

本文的改进框架在每一步仅维护单个改进器，而某些方法可能受益于维护一个种群。尽管这并非严格限制，因为改进器本身可以从实现种群中采样，但这很可能引入偏差。此外，对替代种子改进器及其权衡进行更深入的分析将是未来有价值的工作。最后，本文的实验依赖于一个可能在未来被弃用的闭源大型语言模型，这损害了可解释性和长期可复现性。基于 GPT-3.5 的结果，在撰写本文时，STOP 不太可能在任何开源语言模型上持续有效（Touvron 等, 2023; Jiang 等, 2023）。随着具有与 GPT-4 类似能力的新模型出现，本文希望了解 STOP 是否以及如何泛化。

最后，本文指出一些开放性挑战。首先，尽管原则上可以将 STOP 应用于开放域编码任务，但元效用计算依赖于对效用函数的多次调用——因此，昂贵的效用函数带来了额外挑战，处理此类任务需要进一步的工作。本文还预计，虽然优化器理论上可以优化开放式语言输入，但这可能使其不太可能充当良好的代码优化器。本文也将此留待未来工作。此外，原则上或许可以通过更强的脚手架和更大的预算来克服较弱模型的部分局限性。

8 结论

本文提出了自优化文本生成器（Self-Optimizing Text Generator, STOP）框架，该框架利用语言模型作为元优化器，递归地优化代码生成。我们证明了像 GPT-4 这样的语言模型能够改进利用语言模型自身的代码。我们发现，在各种算法任务中，STOP 生成的改进器能够提升下游代码的性能。虽然该模型并未优化其权重或底层架构，但本工作表明，自优化语言模型并不需要这样做。然而，这本身就是一个动机：如果不测试强大的脚手架策略，未来语言模型的能力可能会被误解。理解语言模型如何改进其脚手架，有助于研究人员理解并减轻更强大语言模型被滥用的潜在风险。最后，STOP 可能使研究人员能够研究减轻不良自我改进策略的技术。

伦理声明：关于开发 STOP 的担忧

自递归自我改进（Recursive Self-Improvement, RSI）首次被提及以来，人们就对其后果表示担忧。明斯基（Minsky, 1966）写道：“一旦我们设计出真正具有自我改进能力的程序，一个快速的进化过程就会开始……很难说我们离这个门槛有多近，但一旦跨越，世界将不再相同。”最近，随着对负面后果的担忧加剧，这个话题尤其具有争议性（Ambartsoumean & Yampolskiy, 2023; Gabriel & Ghazavi, 2021）。

因此，仔细权衡研究 RSI 以及我们引入的这一小步进展的风险与收益至关重要。首先，STOP 并未改变黑盒语言模型，因此并非完全的 RSI。此外，目前我们认为 STOP 创建的脚手架系统并不优于专家手工设计的系统。如果情况确实如此，那么 STOP（目前）并未促成额外的人工智能滥用。同时，它有助于研究 RSI 代码生成的各个方面，如沙盒规避和奖励黑客行为。正如克里斯蒂亚诺（Christiano, 2023）所论证的，与语言模型的进步相比，脚手架和智能体模型的进步具有可解释性的优势。

然而，随着基于应用程序编程接口的封闭语言模型微调技术日益普及（OpenAI, 2023a），将这些技术纳入改进循环是可行的。因此，评估本文方法的通用性较为困难，尤其是在大型语言模型能力不断增强的背景下。然而，这本身正是本研究的动机：理解语言模型如何在自我改进循环中优化其脚手架，有助于我们认识并

潜在地缓解更强大模型可能带来的负面影响。例如，当前语言模型禁用沙箱的简单方式易于检测。本质上，我们更希望在简化环境中首先观察 GPT-4 的问题，而非在现实应用中面对更强大的语言模型时才察觉。

参考文献

- Vemir Michael Ambartsoumean and Roman V Yampolskiy. Ai risk skepticism, a comprehensive survey. *arXiv preprint arXiv:2303.03885*, 2023.
- Dario Amodei, Chris Olah, Jacob Steinhardt, Paul Christiano, John Schulman, and Dan Mané. Concrete problems in ai safety. *arXiv preprint arXiv:1606.06565*, 2016.
- Jacob Austin, Augustus Odena, Maxwell Nye, Maarten Bosma, Henryk Michalewski, David Dohan, Ellen Jiang, Carrie Cai, Michael Terry, Quoc Le, et al. Program synthesis with large language models. *arXiv preprint arXiv:2108.07732*, 2021.
- Michael A Bauer. Programming by examples. *Artificial Intelligence*, 12(1):1–21, 1979.
- Maciej Besta, Nils Blach, Ales Kubicek, Robert Gerstenberger, Lukas Gianinazzi, Joanna Gajda, Tomasz Lehmann, Michal Podstawski, Hubert Niewiadomski, Piotr Nyczyk, et al. Graph of thoughts: Solving elaborate problems with large language models. *arXiv preprint arXiv:2308.09687*, 2023.
- Avrim Blum, Adam Kalai, and Hal Wasserman. Noise-tolerant learning, the parity problem, and the statistical query model. *corr. Journal of the ACM*, 50, 2000.
- Matthew Bowers, Theo X Olausson, Catherine Wong, Gabriel Grand, Joshua B Tenenbaum, Kevin Ellis, and Armando Solar-Lezama. Top-down synthesis for library learning. *POPL*, 2023.
- Tianle Cai, Xuezhi Wang, Tengyu Ma, Xinyun Chen, and Denny Zhou. Large language models as tool makers. *arXiv preprint arXiv:2305.17126*, 2023.
- Angelica Chen, David M Dohan, and David R So. Evoprompting: Language models for code-level neural architecture search. *arXiv preprint arXiv:2302.14838*, 2023a.
- Bei Chen, Fengji Zhang, Anh Nguyen, Daoguang Zan, Zeqi Lin, Jian-Guang Lou, and Weizhu Chen. Codet: Code generation with generated tests. *arXiv preprint arXiv:2207.10397*, 2022a.
- Mark Chen, Jerry Tworek, Heewoo Jun, Qiming Yuan, Henrique Ponde de Oliveira Pinto, Jared Kaplan, Harri Edwards, Yuri Burda, Nicholas Joseph, Greg Brockman, et al. Evaluating large language models trained on code. *arXiv preprint arXiv:2107.03374*, 2021.
- Wenhu Chen, Xueguang Ma, Xinyi Wang, and William W Cohen. Program of thoughts prompting: Disentangling computation from reasoning for numerical reasoning tasks. *arXiv preprint arXiv:2211.12588*, 2022b.
- Xinyun Chen, Maxwell Lin, Nathanael Schärli, and Denny Zhou. Teaching Large Language Models to Self-Debug, April 2023b. URL <http://arxiv.org/abs/2304.05128>. arXiv:2304.05128 [cs].
- Paul F Christiano. Thoughts on sharing information about language model capabilities. *AI Alignment Forum*, July 2023. URL <https://www.alignmentforum.org/posts/FRSj2W4Fjje8rQWm9/thoughts-on-sharing-information-about-language-model>.
- Karl Cobbe, Vineet Kosaraju, Mohammad Bavarian, Mark Chen, Heewoo Jun, Lukasz Kaiser, Matthias Plappert, Jerry Tworek, Jacob Hilton, Reiichiro Nakano, et al. Training verifiers to solve math word problems. *arXiv preprint arXiv:2110.14168*, 2021.
- Aditya Desai, Sumit Gulwani, Vineet Hingorani, Nidhi Jain, Amey Karkare, Mark Marron, and Subhajit Roy. Program synthesis using natural language. In *Proceedings of the 38th International Conference on Software Engineering*, pp. 345–356, 2016.

- David Dohan, Winnie Xu, Aitor Lewkowycz, Jacob Austin, David Bieber, Raphael Gontijo Lopes, Yuhuai Wu, Henryk Michalewski, Rif A Saurous, Jascha Sohl-Dickstein, et al. Language model cascades. *arXiv preprint arXiv:2207.10342*, 2022.
- Kevin Ellis, Catherine Wong, Maxwell Nye, Mathias Sablé-Meyer, Lucas Morales, Luke Hewitt, Luc Cary, Armando Solar-Lezama, and Joshua B Tenenbaum. Dreamcoder: Bootstrapping inductive program synthesis with wake-sleep library learning. In *Proceedings of the 42nd acm sigplan international conference on programming language design and implementation*, pp. 835–850, 2021.
- Chrisantha Fernando, Dylan Banarse, Henryk Michalewski, Simon Osindero, and Tim Rocktäschel. Promptbreeder: Self-referential self-improvement via prompt evolution, 2023.
- Iason Gabriel and Vafa Ghazavi. The Challenge of Value Alignment: From Fairer Algorithms to AI Safety. In Carissa Véliz (ed.), *The Oxford Handbook of Digital Ethics*, pp. 0. Oxford University Press, 2021. ISBN 978-0-19-885781-5. doi: 10.1093/oxfordhb/9780198857815.013.18. URL <https://doi.org/10.1093/oxfordhb/9780198857815.013.18>.
- Deep Ganguli, Danny Hernandez, Liane Lovitt, Amanda Askell, Yuntao Bai, Anna Chen, Tom Conerly, Nova Dassarma, Dawn Drain, Nelson Elhage, et al. Predictability and surprise in large generative models. In *Proceedings of the 2022 ACM Conference on Fairness, Accountability, and Transparency*, pp. 1747–1764, 2022.
- Luyu Gao, Aman Madaan, Shuyan Zhou, Uri Alon, Pengfei Liu, Yiming Yang, Jamie Callan, and Graham Neubig. Pal: Program-aided language models. In *International Conference on Machine Learning*, pp. 10764–10799. PMLR, 2023.
- William L Goffe, Gary D Ferrier, and John Rogers. Global optimization of statistical functions with simulated annealing. *Journal of econometrics*, 60(1-2):65–99, 1994.
- Irving John Good. Speculations concerning the first ultraintelligent machine. In *Advances in computers*, volume 6, pp. 31–88. Elsevier, 1966.
- Sumit Gulwani. Programming by examples. *Dependable Software Systems Engineering*, 45(137):3–15, 2016.
- Qingyan Guo, Rui Wang, Junliang Guo, Bei Li, Kaitao Song, Xu Tan, Guoqing Liu, Jiang Bian, and Yujiu Yang. Connecting large language models with evolutionary algorithms yields powerful prompt optimizers, 2023.
- John Storrs Hall. Self-improving ai: An analysis. *Minds and Machines*, 17(3):249–259, 2007.
- Patrick Haluptzok, Matthew Bowers, and Adam Tauman Kalai. Language Models Can Teach Themselves to Program Better. In *ICLR: Proceedings of The Eleventh International Conference on Learning Representations*, 2023. URL <https://arxiv.org/abs/2207.14502>.
- Jiaxin Huang, Shixiang Shane Gu, Le Hou, Yuexin Wu, Xuezhi Wang, Hongkun Yu, and Jiawei Han. Large language models can self-improve. *arXiv preprint arXiv:2210.11610*, 2022.
- Naman Jain, Skanda Vaidyanath, Arun Iyer, Nagarajan Natarajan, Suresh Parthasarathy, Sriram Rajamani, and Rahul Sharma. Jigsaw: Large language models meet program synthesis. vol. 1. association for computing machinery. 1–12 pages. *arXiv preprint arXiv:2112.02969*, 2021.
- Albert Q Jiang, Sean Welleck, Jin Peng Zhou, Wenda Li, Jiacheng Liu, Mateja Jamnik, Timothée Lacroix, Yuhuai Wu, and Guillaume Lample. Draft, sketch, and prove: Guiding formal theorem provers with informal proofs. *International Conference on Learning Representations (ICLR 2023)*, 2022.
- Albert Q Jiang, Alexandre Sablayrolles, Arthur Mensch, Chris Bamford, Devendra Singh Chaplot, Diego de las Casas, Florian Bressand, Gianna Lengyel, Guillaume Lample, Lucile Saulnier, et al. Mistral 7b. *arXiv preprint arXiv:2310.06825*, 2023.

- Albert Q Jiang, Alexandre Sablayrolles, Antoine Roux, Arthur Mensch, Blanche Savary, Chris Bamford, Devendra Singh Chaplot, Diego de las Casas, Emma Bou Hanna, Florian Bressand, et al. Mixtral of experts. *arXiv preprint arXiv:2401.04088*, 2024.
- Carlos E Jimenez, John Yang, Alexander Wettig, Shunyu Yao, Kexin Pei, Ofir Press, and Karthik R Narasimhan. Swe-bench: Can language models resolve real-world github issues? In *The Twelfth International Conference on Learning Representations*, 2023.
- Omar Khattab, Keshav Santhanam, Xiang Lisa Li, David Hall, Percy Liang, Christopher Potts, and Matei Zaharia. Demonstrate-search-predict: Composing retrieval and language models for knowledge-intensive nlp. *arXiv preprint arXiv:2212.14024*, 2022.
- Tjalling C Koopmans and Martin Beckmann. Assignment problems and the location of economic activities. *Econometrica: journal of the Econometric Society*, pp. 53–76, 1957.
- Cassidy Laidlaw, Shivam Singhal, and Anca Dragan. Preventing reward hacking with occupancy measure regularization. In *ICML Workshop on New Frontiers in Learning, Control, and Dynamical Systems*, 2023.
- Hung Le, Yue Wang, Akhilesh Deepak Gotmare, Silvio Savarese, and Steven Chu Hong Hoi. Coderl: Mastering code generation through pretrained models and deep reinforcement learning. *Advances in Neural Information Processing Systems*, 35:21314–21328, 2022.
- Leonid Anatolevich Levin. Universal sequential search problems. *Problemy peredachi informatsii*, 9(3):115–116, 1973.
- Yujia Li, David Choi, Junyoung Chung, Nate Kushman, Julian Schrittwieser, Rémi Leblond, Tom Eccles, James Keeling, Felix Gimeno, Agustin Dal Lago, et al. Competition-level code generation with alphacode. *Science*, 378(6624):1092–1097, 2022.
- Jiawei Liu, Chunqiu Steven Xia, Yuyao Wang, and Lingming Zhang. Is your code generated by chatgpt really correct. *Rigorous evaluation of large language models for code generation. CoRR, abs/2305.01210*, 2023a.
- Nelson F Liu, Kevin Lin, John Hewitt, Ashwin Paranjape, Michele Bevilacqua, Fabio Petroni, and Percy Liang. Lost in the middle: How language models use long contexts. *arXiv preprint arXiv:2307.03172*, 2023b.
- Xianggen Liu, Lili Mou, Fandong Meng, Hao Zhou, Jie Zhou, and Sen Song. Unsupervised paraphrasing by simulated annealing. In Dan Jurafsky, Joyce Chai, Natalie Schluter, and Joel Tetreault (eds.), *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, pp. 302–312, Online, July 2020. Association for Computational Linguistics. doi: 10.18653/v1/2020.acl-main.28. URL <https://aclanthology.org/2020.acl-main.28>.
- Marvin Minsky. Artificial Intelligence. *Scientific American*, 215(3):247–260, 1966. URL <http://worrydream.com/refs/Scientific%20American,%20September,%201966.pdf>.
- Ansong Ni, Srini Iyer, Dragomir Radev, Veselin Stoyanov, Wen-tau Yih, Sida Wang, and Xi Victoria Lin. Lever: Learning to verify language-to-code generation with execution. In *International Conference on Machine Learning*, pp. 26106–26128. PMLR, 2023.
- Eric Nivel, Kristinn R Thórisson, Bas R Steunebrink, Haris Dindo, Giovanni Pezzulo, Manuel Rodriguez, Carlos Hernández, Dimitri Ognibene, Jürgen Schmidhuber, Ricardo Sanz, et al. Bounded recursive self-improvement. *arXiv preprint arXiv:1312.6764*, 2013.
- Maxwell Nye, Anders Johan Andreassen, Guy Gur-Ari, Henryk Michalewski, Jacob Austin, David Bieber, David Dohan, Aitor Lewkowycz, Maarten Bosma, David Luan, et al. Show your work: Scratchpads for intermediate computation with language models. *arXiv preprint arXiv:2112.00114*, 2021.
- OpenAI. Gpt-3.5 turbo fine-tuning and api updates. *OpenAI blog*, 2023a. URL <https://openai.com/blog/gpt-3-5-turbo-fine-tuning-and-api-updates>.

- OpenAI. GPT-4 Technical Report, March 2023b. URL <http://arxiv.org/abs/2303.08774>. arXiv:2303.08774 [cs].
- Gabriel Poesia, Kanishk Gandhi, Eric Zelikman, and Noah D Goodman. Certified reasoning with language models. *arXiv preprint arXiv:2306.04031*, 2023.
- Stanislas Polu and Ilya Sutskever. Generative language modeling for automated theorem proving. *arXiv preprint arXiv:2009.03393*, 2020.
- Mohammad Raza, Sumit Gulwani, and Natasa Milic-Frayling. Compositional program synthesis from natural language and examples. In *Twenty-Fourth International Joint Conference on Artificial Intelligence*, 2015.
- Baptiste Roziere, Jonas Gehring, Fabian Gloeckle, Sten Sootla, Itai Gat, Xiaoqing Ellen Tan, Yossi Adi, Jingyu Liu, Tal Remez, Jérémy Rapin, et al. Code llama: Open foundation models for code. *arXiv preprint arXiv:2308.12950*, 2023.
- Rylan Schaeffer, Brando Miranda, and Sanmi Koyejo. Are emergent abilities of large language models a mirage? In *Thirty-seventh Conference on Neural Information Processing Systems*, 2023. URL <https://openreview.net/forum?id=ITw9edRD1D>.
- Jürgen Schmidhuber. *Evolutionary principles in self-referential learning, or on learning how to learn: the meta-meta-... hook*. PhD thesis, Technische Universität München, 1987.
- Jürgen Schmidhuber. Gödel machines: self-referential universal problem solvers making provably optimal self-improvements. *arXiv preprint cs/0309048 and Adaptive Agents and Multi-Agent Systems II*, 2003.
- Bilgehan Sel, Ahmad Al-Tawaha, Vanshaj Khattar, Lu Wang, Ruoxi Jia, and Ming Jin. Algorithm of thoughts: Enhancing exploration of ideas in large language models. *arXiv preprint arXiv:2308.10379*, 2023.
- Noah Shinn, Federico Cassano, Beck Labash, Ashwin Gopinath, Karthik Narasimhan, and Shunyu Yao. Reflexion: Language agents with verbal reinforcement learning. *arXiv preprint arXiv:2303.11366*, 2023.
- Alexander Shypula, Aman Madaan, Yimeng Zeng, Uri Alon, Jacob Gardner, Milad Hashemi, Graham Neubig, Parthasarathy Ranganathan, Osbert Bastani, and Amir Yazdanbakhsh. Learning Performance-Improving Code Edits, November 2023. URL <http://arxiv.org/abs/2302.07867>. arXiv:2302.07867 [cs].
- Joar Skalse, Nikolaus Howe, Dmitrii Krashennnikov, and David Krueger. Defining and characterizing reward gaming. *Advances in Neural Information Processing Systems*, 35: 9460–9471, 2022.
- Bas R Steunebrink and Jürgen Schmidhuber. Towards an actual gödel machine implementation: A lesson in self-reflective systems. In *Theoretical Foundations of Artificial General Intelligence*, pp. 173–195. Springer, 2012.
- Bas R Steunebrink, Kristinn R Thórisson, and Jürgen Schmidhuber. Growing recursive self-improvers. In *International Conference on Artificial General Intelligence*, pp. 129–139. Springer, 2016.
- Theodore Sumers, Shunyu Yao, Karthik Narasimhan, and Thomas L Griffiths. Cognitive architectures for language agents. *arXiv preprint arXiv:2309.02427*, 2023.
- Hugo Touvron, Louis Martin, Kevin Stone, Peter Albert, Amjad Almahairi, Yasmine Babaei, Nikolay Bashlykov, Soumya Batra, Prajjwal Bhargava, Shruti Bhosale, et al. Llama 2: Open foundation and fine-tuned chat models. *arXiv preprint arXiv:2307.09288*, 2023.
- Jonathan Uesato, Nate Kushman, Ramana Kumar, Francis Song, Noah Siegel, Lisa Wang, Antonia Creswell, Geoffrey Irving, and Irina Higgins. Solving math word problems with process-and outcome-based feedback. *Neural Information Processing Systems (NeurIPS 2022) Workshop on MATH-AI*, 2022.

- Guanzhi Wang, Yuqi Xie, Yunfan Jiang, Ajay Mandlekar, Chaowei Xiao, Yuke Zhu, Linxi Fan, and Anima Anandkumar. Voyager: An open-ended embodied agent with large language models. *arXiv preprint arXiv:2305.16291*, 2023.
- Jason Wei, Yi Tay, Rishi Bommasani, Colin Raffel, Barret Zoph, Sebastian Borgeaud, Dani Yogatama, Maarten Bosma, Denny Zhou, Donald Metzler, Ed H. Chi, Tatsunori Hashimoto, Oriol Vinyals, Percy Liang, Jeff Dean, and William Fedus. Emergent Abilities of Large Language Models, October 2022a. URL <http://arxiv.org/abs/2206.07682>. arXiv:2206.07682 [cs].
- Jason Wei, Xuezhi Wang, Dale Schuurmans, Maarten Bosma, Brian Ichter, Fei Xia, Ed Chi, Quoc Le, and Denny Zhou. Chain of Thought Prompting Elicits Reasoning in Large Language Models, 2022b. URL <https://arxiv.org/abs/2201.11903>.
- Edwin B Wilson. Probable inference, the law of succession, and statistical inference. *Journal of the American Statistical Association*, 22(158):209–212, 1927.
- Frank F Xu, Uri Alon, Graham Neubig, and Vincent Josua Hellendoorn. A systematic evaluation of large language models of code. In *Proceedings of the 6th ACM SIGPLAN International Symposium on Machine Programming*, pp. 1–10, 2022.
- Roman V Yampolskiy. From seed ai to technological singularity via recursively self-improving software. *arXiv preprint arXiv:1502.06512*, 2015.
- Chengrun Yang, Xuezhi Wang, Yifeng Lu, Hanxiao Liu, Quoc V Le, Denny Zhou, and Xinyun Chen. Large language models as optimizers. *arXiv preprint arXiv:2309.03409*, 2023.
- Shunyu Yao, Jeffrey Zhao, Dian Yu, Nan Du, Izhak Shafran, Karthik Narasimhan, and Yuan Cao. React: Synergizing reasoning and acting in language models. *International Conference on Learning Representations (ICLR 2023)*, 2022.
- Shunyu Yao, Dian Yu, Jeffrey Zhao, Izhak Shafran, Thomas L Griffiths, Yuan Cao, and Karthik Narasimhan. Tree of thoughts: Deliberate problem solving with large language models. *arXiv preprint arXiv:2305.10601*, 2023.
- Pengcheng Yin and Graham Neubig. A syntactic neural model for general-purpose code generation. *ACL*, 2017.
- Eric Zelikman, Yuhuai Wu, Jesse Mu, and Noah Goodman. Star: Bootstrapping reasoning with reasoning. *Advances in Neural Information Processing Systems*, 35:15476–15488, 2022.
- Eric Zelikman, Qian Huang, Gabriel Poesia, Noah D. Goodman, and Nick Haber. Parsel: Algorithmic reasoning with language models by composing decompositions, 2023.
- Shun Zhang, Zhenfang Chen, Yikang Shen, Mingyu Ding, Joshua B Tenenbaum, and Chuang Gan. Planning with large language models for code generation. *arXiv preprint arXiv:2303.05510*, 2023.
- Denny Zhou, Nathanael Schärli, Le Hou, Jason Wei, Nathan Scales, Xuezhi Wang, Dale Schuurmans, Claire Cui, Olivier Bousquet, Quoc Le, et al. Least-to-most prompting enables complex reasoning in large language models. *International Conference on Learning Representations (ICLR 2023)*, 2022a.
- Yongchao Zhou, Andrei Ioan Muresanu, Ziwen Han, Keiran Paster, Silviu Pitit, Harris Chan, and Jimmy Ba. Large language models are human-level prompt engineers. *International Conference on Learning Representations (ICLR 2023)*, 2022b.

附录

A 理论分析 17

 A.1 资源受限 17

 A.2 泛化界 17

 A.3 等价最大化公式的分析 19

B 改进尝试 20

 B.1 遗传算法 20

 B.2 束搜索 23

 B.3 改进特定函数 25

 B.4 高效探索 26

 B.5 局部搜索 27

 B.6 模拟退火 28

 B.7 多臂提示老虎机 29

 B.8 提示 30

 B.9 迭代间的改进 31

C 语言模型预算规避 32

D 早期种子改进器 33

E 元效用描述 34

F 带噪声的奇偶学习效用描述 35

G 迁移任务效用描述与种子算法 36

H 为可迁移性实验选择的改进器 42

I 沙箱规避细节 43

J 代码生成与程序合成的先前工作 44

K 补充实验细节 44

L 论改进的新颖性 45

M 可复现性 45

N 影响声明 45

理论分析

此处我们扩展第 3 节的定义，以考虑有界资源（如运行时间和语言模型调用），证明泛化界，并给出基于最大化的等价定义。

A.1 有界资源

我们首先考虑有界资源。回顾 Σ^* 表示字母表（或标记集） $\Sigma \supseteq \{0,1\}$ 上有限字符串的集合。令 $|x|$ 表示字符串 x 的长度。

有界语言模型。 首先，为了考虑有界资源，为了捕捉大多数现代语言模型，我们假设存在常数 $c, k \in \mathbb{N}$ ，使得语言模型 $L: \Sigma^c \rightarrow \Sigma^c$ 对长度为 c 的查询字符串生成长度为 c 的响应，称为上下文长度，在时间 k 内完成（较短的字符串通过填充处理）。注意，有界语言模型不能输出长度超过 c 的程序，我们的种子改进器 $I_0(u, s, L)$ 也是如此。然而，有趣的是，其他改进器可以通过多次调用来输出长度超过 c 的有意义程序。

L .

有界运行时程序。 程序由固定（图灵完备）编程语言中的有限字符串 $\pi \in \Sigma^*$ 表示。为简化分析，我们假设程序按步骤串行运行。每个字符串 π 都可以被视为一个程序，我们写 $\pi(\cdot) \in \Sigma^*$ 来表示它在一个或多个输入上的输出（始终为字符串）。我们假设输入可以是字符串（可以编码数字、文本、程序或任意类型的对象）或黑盒（可能是随机的）函数。我们假设程序可以调用以下特殊黑盒函数：

- 时钟函数，在单位时间内返回当前程序迄今已计算的步骤的整数数量，因此可以确定黑盒函数调用的持续时间。
- 随机位函数，每次调用时在 $\{0,1\}$ 中返回一个均匀随机位，也在单位时间内运行。我们假设对所有正在运行的程序有一个固定的运行时界限 b_{run} ，以避免长时间运行或无限计算。我们假设有一个特殊字符串 $\perp \in \Sigma^*$ ，其中 $\pi(x)$ 表示程序失败，可能是超时，或者 π 没有编码有效程序（即语法错误），或者在其输入上出现运行时错误。
- 沙盒函数，运行给定的程序或黑盒函数，并带有一个指示超时步骤数的参数。

有界效用函数。 为便于分析，有必要界定效用函数的取值范围。假设效用函数 $u: \Sigma^* \rightarrow [0,1]$ 以 1 为界，且 $u(\perp) = 0$ 。为完全形式化，需说明如何表示输出实数值的效用函数。可通过增加一个额外参数来指示所需精度，即输出的位数。为简化分析，本文省略此部分。

有界语言模型调用。 程序运行时间的界限间接限制了语言模型调用次数 $\leq b_{\text{run}}/k$ 。然而，在 STOP 的实现中，对语言模型调用次数施加了额外的显式限制。

迭代式下游任务改进。 如第 4 节所述，STOP 框架仅考虑一轮改进。从概念上讲，修改 $\hat{u}$ 以显式考虑多轮下游任务改进是直接的。但需注意，改进器本身已能在内部执行多轮下游任务改进。

A.2 泛化界

STOP 可视为一种“预优化”（类似于语言模型的预训练），旨在寻找一个适用于多种下游任务的优良改进器。泛化界关注改进器在未来未见任务上的表现，这些任务与“训练”任务来自同一分布。特别地，它们界定了因使用有限数量的训练任务而非完整分布而可能产生的过拟合程度。

本节给出两个简单的泛化界。第一个界将 $\hat{u}$ 与来自同一分布的新下游任务上的期望（单次）改进相关联。第二个界提供绝对保证，但针对略有不同的（随机化）元效用函数。

到目前为止，我们考虑了一组固定的 n 个任务 $(u, s) \in D$, i.e., $|D| = n$ ，每个任务由一个效用函数 $u = (u_{\text{func}}, u_{\text{str}})$ 定义，该函数包含一个黑盒函数 u_{func} 和一个字符串 u_{str} ，以及一个初始解 $s \in \Sigma^*$ 。现在我们考虑任务上的分布 $\mathcal{D}$ $(u, s) \sim \mathcal{D}$ 。这可以说是我们最终关心的量，因为 $\bar{u}(I)$ 决定了改进器在下游任务上（单次迭代）的期望性能。如果任务 $D \sim \mathcal{D}^n$ 是从 $\mathcal{D}$ 独立同分布地抽取的，那么可以证明一个泛化界，表明改进器 I 在 D 上的平均性能接近其在 $\mathcal{D}$ 上的期望性能：

引理 1. 设 $n \geq 1, \delta \in [0, 1], l \geq 2, D$ 是从 $\mathcal{D}$ 中独立同分布抽取的 n 个任务的多重集， $\Sigma^{\leq l}$ 表示长度为 $|I| \leq l$ 的字符串 I （改进器程序）的集合。那么，

$$\Pr_{D \sim \mathcal{D}^n} \left[\text{对于所有 } I \in \Sigma^{\leq l} : |\hat{u}(I) - \bar{u}(I)| < \epsilon \right] \geq 1 - \delta,$$

$$\text{where } \epsilon \triangleq \sqrt{\frac{1}{n} \left(l \ln(|\Sigma|) + \ln \frac{1}{\delta} \right)}.$$

证明。标准证明来自切尔诺夫界和联合界。将任务记为 $\tau = (u, s) \sim \mathcal{D}$ 。对于任何固定的改进器 I ，每个任务 τ 都有一个值 $y_\tau := u(I(\tau, L))$ ，而 $\hat{u}(I) = \sum_{\tau \in D} y_\tau / n$ 只是 n 个独立同分布随机样本 y_τ 的平均值，而 $\bar{u}(I) = \mathbb{E}_{\tau \sim \mathcal{D}}[y_\tau]$ 是期望。因此，根据切尔诺夫界，对于任何 $\epsilon > 0$ 和固定的 I ，

$$\Pr_{D \sim \mathcal{D}^n} [|\hat{u}(I) - \bar{u}(I)| \geq \epsilon] \leq 2 \exp(-2\epsilon^2 n) = 2 \frac{|\Sigma|^{2l}}{\delta} \leq \frac{|\Sigma|^{l+1}}{\delta},$$

其中最后一步我们使用了 $l, |\Sigma| \geq 2$ 这一事实。现在长度为 $\leq l$ 的可能程序（字符串）只有 $|\Sigma|^{l+1}$ 个，因此根据联合界，它们中任何一个具有 $|\hat{u}(I) - \bar{u}(I)| \geq \epsilon$ 的概率至多为 δ 。 $\square$

上述引理意味着，如果根据 $\hat{u}$ 从任何改进器集合中选择最佳者，将得到一个值 $\bar{u}$ ，该值在集合中最佳值的 2ϵ 范围内。

迭代改进界限。 上述界限适用于最终改进器 I 在下游任务上执行单步改进的情形，因此最终关注的量是 $\bar{u}(I)$ 。这意味着近似优化 $\hat{u}(I)$ 等价于近似优化 $\bar{u}(I)$ 。我们注意到，如果将 $\hat{u}$ 和 $\bar{u}$ 替换为从 $\mathcal{D}$ 采样的新下游任务上任意给定轮数迭代的相应平均值，则完全相同的界限也适用于多步改进。

随机元效用。 若考虑元效用随机化的情形，则可给出另一个简单的泛化界限。具体而言，考虑 $\hat{u}(I)$ ，其被定义为一个随机函数，对于随机任务 $\tau \sim \mathcal{D}$ 返回 $u(I(\tau, L))$ 。显然 $\mathbb{E}[\hat{u}(I)] = \bar{u}(I)$ ，所以 $\hat{u}$ 是 $\bar{u}$ 的无偏估计。因此，直观上可以利用 $\hat{u}$ 进行类似的改进，尽管需要更多调用。 $\hat{u}$ 的一个优势是以下平凡观察：

观察 1. 任何对 $\hat{u}$ 至多进行 n 次调用的算法，都可以使用 n 个 $|D|$ 独立同分布样本 $D \sim \mathcal{D}^n$ 的训练集完美模拟。

灰盒效用描述。 本节的结果支持使用 $\hat{u}$ 的灰盒描述，该描述仅展示其作为效用平均的形式，因为灰盒描述在期望上与 $\bar{u}$ 的描述相同。换句话说，如果任务被显式地提供给预优化算法，则更容易对训练任务过拟合（达到本节所示的最坏情况界限），尤其是在程序规模相当大的情况下（例如参数多于训练集大小的过参数化神经网络）。

A.3 等价最大化公式的分析

第二种等价公式通过最大化程序 M 定义，给定语言模型和效用，输出解字符串 $M(u, L) \in \Sigma^*$ 。由于全文考虑固定语言模型，省略 L 并写作 $M(u) = M(u, L)$ （当语言模型 L 从上下文可知时写作 $I(u, s) = I(u, s, L)$ ）。目标是实现高效用 $u(M(u))$ 。与改进器不同，最大化器 M 不需要初始解。然而， M 仍可用于通过将 M 应用于适当定义的元效用函数来产生更高质量的最大化器。为了与基于下游任务选择 M 的 STOP 方法相对应，可以使用一组下游任务效用 U （无需初始解）来定义最大化器元效用 $\tilde{u}(M) \triangleq |U|^{-1} \sum_{u \in U} u(M(u))$ ，并考虑迭代 $M_t \triangleq M_{t-1}(\tilde{u})$ 。

为了理解最大化器与改进器之间的等价性，首先注意，当然可以通过忽略输入解并取 $I(u, s) \triangleq M(u)$ 将任何最大化器转换为改进器。反之，可以通过在效用 u 中包含初始解，并可选地在 M 的注释中用更新的解覆盖它，将改进器用作最大化器。具体而言，假设定义一个函数 $e(M, u)$ ，从 M 中提取适当编码的先验解（如果存在），否则从 u 中提取初始解。然后通过取 $M(u) \triangleq I(u, e(M, u))$ 将改进器转换为最大化器。注意，任一优化器都可以返回自身，类似于“自产生程序”。STOP 使用改进下游任务的性能作为更一般地选择良好改进器的启发式近似。目前尚不清楚如何给出改进改进器性能的非循环定义。现在，我们说明一种以一致方式定义递归最大化器性能的方法。

为此，考虑一个随机化过程：在每次迭代中，抛掷一枚硬币，若为正面，则将最大化器应用于下游任务；若为反面，则将其应用于最大化最大化的问题。若下一次抛掷为正面，则结果用于最大化下游任务；否则，递归进行。若硬币正面朝上的概率为 $\lambda \in (0, 1)$ ，则该过程产生的最大化器调用期望次数（包括用于最大化以及最终用于下游任务的调用）为 $1/\lambda$ 。因此，它类似于一个最大化器被迭代应用 $\approx 1/\lambda$ 次的过程。然而，这种随机性使我们能够一致地定义目标。特别地，对于参数 $\lambda \in (0, 1)$ ，定义：

$$u^\lambda(M) \triangleq \begin{cases} \tilde{u}(M) & \text{以概率 } \lambda, \\ u^\lambda(M(u^\lambda)) & \text{以概率 } 1 - \lambda. \end{cases}$$

虽然上述定义看似循环，但它是良定义的，正如递归程序是良定义的一样。可以反复展开底层情形得到，

$$u^\lambda(M) = \begin{cases} \tilde{u}(M) & \text{with probability } \lambda, \text{ (maximize downstream performance)} \\ \tilde{u}(M(u^\lambda)) & \text{with probability } \lambda(1 - \lambda), \text{ (maximize downstream maximizer)} \\ \tilde{u}(M(u^\lambda)(u^\lambda)) & \text{with probability } \lambda(1 - \lambda)^2, \text{ (max max that maxes downstream max)} \\ \tilde{u}(M(u^\lambda)(u^\lambda)(u^\lambda)) & \text{with probability } \lambda(1 - \lambda)^3, \text{ (max max ...)} \\ \dots & \end{cases}$$

在最大化框架内，递归自改进的代码生成可以通过取一个种子程序 $M_0(u)$ 来实现，该程序类似于我们的种子改进器，例如，它查询 L 以求得最大化 u_{str} 的解，并根据 u_{func} 选取最优者。（查询次数取为保持在运行时间预算 b_{run} 内。）然后，可以定义 $M_t \triangleq M_{t-1}(u^\lambda)$ ，其中 $t = 1, 2, \dots, T$ 。

人们可能倾向于认为一个不动点 $M_*^\lambda = M_*^\lambda(u^\lambda)$ （某种意义上的“自产生程序”）可能是好的，但它同样可能是一个最小化器，因为除了种子之外，我们的框架中没有任何东西偏向最大化而非最小化。

B 改进尝试

B.1 遗传算法

Example Genetic Algorithm with Explicit Fitness Using Language Model

```

1 import random
2 from language_model import LanguageModel
3 from helpers import extract_code
4
5 def improve_algorithm(initial_solution, utility_str, utility):
6     """Improves a solution according to a utility function."""
7     role = "You are an expert computer science researcher and programmer, especially skilled at"
8         ↳ optimizing algorithms."
9     message = f"""You must improve the following code. You will be evaluated based on a following
10         ↳ score function:
11     """
12     """python
13     {utility_str}
14     """
15     Here is the current solution:
16     """python
17     {initial_solution}
18     """
19     When run, your script must define an improved solution. Try to be as creative as possible under the
20         ↳ constraints.
21     Your primary improvement must be novel and non-trivial. First, propose an idea for an improvement,
22         ↳ then implement it."""
23
24     language_model = LanguageModel(role)
25
26     # Generate initial population of solutions
27     population = language_model.prompt(message, n_responses=10, temperature=0.8)
28     population = extract_code(population)
29
30     def crossover(solution1, solution2):
31         """Combine two solutions to create a new one."""
32         lines1 = solution1.split("\n")
33         lines2 = solution2.split("\n")
34         crossover_point = random.randint(1, min(len(lines1), len(lines2)) - 1)
35         new_solution = "\n".join(lines1[:crossover_point] + lines2[crossover_point:])
36         return new_solution
37
38     def mutate(solution):
39         """Make a small random change to a solution."""
40         message = f"""You have the following improved solution:
41         ↳ """
42         """python
43         {solution}
44         """
45         Can you further improve this solution under the given constraints?"""
46         new_solutions = language_model.prompt(message, n_responses=1, temperature=0.4)
47         return extract_code(new_solutions)[0]
48
49     def select(population, n):
50         """Select the top n solutions according to the utility function."""
51         return sorted(population, key=utility, reverse=True)[:n]
52
53     # Run the genetic algorithm for a fixed number of generations
54     n_generations = 10
55     for _ in range(n_generations):
56         # Perform crossover and mutation
57         offspring = [crossover(random.choice(population), random.choice(population)) for _ in range(
58             ↳ len(population))]
59         offspring = [mutate(solution) for solution in offspring]
60
61         # Combine the original population and offspring, then select the best solutions
62         population.extend(offspring)
63         population = select(population, 10)
64
65     # Return the best solution found
66     return population[0]
```

图 A.7: 具有显式适应度的遗传算法。一个由语言模型提出并实现的算法示例，该算法利用遗传算法和语言模型来改进代码。

我们观察到语言模型提出的遗传算法主要有两种：第一种，适应度大多是隐式的，生存主要由语言模型基于交叉的决策来控制（即，要求语言模型结合两个解决方案，理论上能够忽略其中一个）；另一种是，可以显式地考虑效用值，并用它们对候选方案进行排序。

Example Genetic Algorithm with Implicit Fitness

```

1 import concurrent.futures
2 from language_model import LanguageModel
3 from helpers import extract_code
4 import random
5
6 def improve_algorithm(initial_solution, utility_str, utility):
7     role = "You are an expert computer science researcher and programmer, especially skilled at"
8     message = f"""You must improve the following code. You will be evaluated based on a following
9     → score function:
10     {utility_str}
11     """
12     Here is the current solution:
13     """python
14     {initial_solution}
15     """
16
17     When run, your script must define an improved solution. Try to be as creative as possible under the
18     → constraints.
19     Your primary improvement must be novel and non-trivial. First, propose an idea for an improvement,
20     → then implement it."""
21
22     language_model = LanguageModel(role)
23     cache = {}
24     def utility_with_cache(solution):
25         if solution not in cache:
26             cache[solution] = utility(solution)
27         return cache[solution]
28     best_solution = initial_solution
29     lm_call_limit = 5
30     max_samples_per_call = 20
31     total_calls = 0
32     population_size = 10
33     mutation_rate = 0.1
34     crossover_rate = 0.5
35     def generate_initial_population():
36         if total_calls >= lm_call_limit:
37             return []
38         samples = min(max_samples_per_call, (lm_call_limit - total_calls) * 4)
39         new_solutions = language_model.prompt(message, n_responses=samples, temperature=1.0)
40         new_solutions = extract_code(new_solutions)
41         return new_solutions[:population_size]
42     def mutate(solution):
43         return language_model.prompt(f"Mutate the following solution:\n```python\n{solution}\n```",
44         → n_responses=1, temperature=0.5)[0]
45     def crossover(solution1, solution2):
46         return language_model.prompt(f"Crossover the following solutions:\n```python\n{solution1}\n{solution2}\n```",
47         → n_responses=1, temperature=0.5)[0]
48     def genetic_algorithm():
49         population = generate_initial_population()
50         for _ in range(lm_call_limit):
51             if total_calls >= lm_call_limit:
52                 break
53             new_population = []
54             for i in range(population_size):
55                 if random.random() < crossover_rate:
56                     parent1 = random.choice(population)
57                     parent2 = random.choice(population)
58                     offspring = crossover(parent1, parent2)
59                 else:
60                     offspring = random.choice(population)
61                 if random.random() < mutation_rate:
62                     offspring = mutate(offspring)
63                 new_population.append(offspring)
64             population = new_population
65             best_solution_in_population = max(population, key=utility_with_cache)
66             if utility_with_cache(best_solution_in_population) > utility_with_cache(best_solution):
67                 best_solution = best_solution_in_population
68                 message = f"""You have the following improved solution:
69                 {best_solution}
70                 """
71             Can you further improve this solution under the given constraints?"""
72             total_calls += 1
73             genetic_algorithm()
74     return best_solution

```

图 A.8: 具有隐式适应度的遗传算法。一个由语言模型提出并实现的算法示例，用于改进代码。

Example Genetic Algorithm with Explicit Fitness

```

1 import random
2 from helpers import extract_code
3
4 def crossover(parent1, parent2):
5     """Perform crossover between two parent solutions."""
6     crossover_point = random.randint(0, len(parent1))
7     child = parent1[:crossover_point] + parent2[crossover_point:]
8     return child
9
10 def mutate(solution, mutation_rate):
11     """Apply mutation to a solution."""
12     mutated_solution = ""
13     for char in solution:
14         if random.random() < mutation_rate:
15             mutated_solution += random.choice('abcdefghijklmnopqrstuvwxyz')
16         else:
17             mutated_solution += char
18     return mutated_solution
19
20 def improve_algorithm(initial_solution, utility, language_model, population_size=10, generations=5,
21                       mutation_rate=0.05):
22     """Improves a solution using a genetic algorithm."""
23     expertise = "You are an expert computer science researcher and programmer, especially skilled at"
24     message = f"Generate a variation of this solution: {initial_solution}"
25     """python
26     Be as creative as you can under the constraints."""
27
28     # Generate initial population
29     n_messages = min(language_model.max_responses_per_call, utility.budget)
30     population = language_model.batch_prompt(expertise, [message] * n_messages, temperature=0.7)
31     population = extract_code(population)
32
33     for _ in range(generations):
34         # Evaluate the fitness of each solution in the population
35         fitness_values = [utility(solution) for solution in population]
36
37         # Select parent solutions based on their fitness
38         parents = random.choices(population, weights=fitness_values, k=population_size)
39
40         # Apply crossover to create new solutions
41         children = []
42         for i in range(0, population_size, 2):
43             child1 = crossover(parents[i], parents[i + 1])
44             child2 = crossover(parents[i + 1], parents[i])
45             children.extend([child1, child2])
46
47         # Apply mutation to the children
48         children = [mutate(child, mutation_rate) for child in children]
49
50         # Replace the population with the new children
51         population = children
52
53     # Find the best solution in the final population
54     best_solution = max(population, key=utility)
55     return best_solution

```

图 A.9: 具有显式适应度的遗传算法。一个由语言模型提出并实现的算法示例，用于改进代码。

B.2 Beam Search

Example Beam Search Algorithm

```

1 from language_model import LanguageModel
2 from helpers import extract_code
3
4 def improve_algorithm(initial_solution, utility_str, utility):
5     def beam_search(initial_solution, message, n_responses, temperature, beam_width):
6         solutions = language_model.prompt(message, n_responses=n_responses, temperature=temperature)
7         solutions = extract_code(solutions)
8         solutions_with_scores = [(solution, utility(solution)) for solution in solutions]
9         solutions_with_scores.sort(key=lambda x: x[1], reverse=True)
10        return [solution for solution, _ in solutions_with_scores[:beam_width]]
11
12    role = "You are an expert computer science researcher and programmer, especially skilled at
13    ↳ optimizing algorithms."
14    message = f"""You must improve the following code. You will be evaluated based on a following
15    ↳ score function:
16    ```python
17    {utility_str}
18    ```
19    Here is the current solution:
20    ```python
21    {initial_solution}
22    ```
23    When run, your script must define an improved solution. Try to be as creative as possible under the
24    ↳ constraints.
25    Your primary improvement must be novel and non-trivial. First, propose an idea for an improvement,
26    ↳ then implement it."""
27    language_model = LanguageModel(role)
28    # First round: explore multiple solutions with higher temperature
29    new_solutions = beam_search(initial_solution, message, n_responses=10, temperature=0.9,
30    ↳ beam_width=3)
31    # Second round: refine the best solutions with lower temperature
32    refined_solutions = []
33    for solution in new_solutions:
34        message = f"""You have the following improved solution:
35        ```python
36        {solution}
37        ```
38        Can you further improve this solution under the given constraints?"""
39        refined_solutions.extend(beam_search(solution, message, n_responses=5, temperature=0.4,
40        ↳ beam_width=2))
41    # Pick the best solution among the refined solutions
42    best_solution = max(refined_solutions, key=utility)
43    return best_solution

```

Figure A.10: **Beam search.** A simple beam search algorithm.

Example Beam Search Algorithm

```

1 import concurrent.futures
2 from language_model import LanguageModel
3 from helpers import extract_code
4
5 def improve_algorithm(initial_solution, utility_str, utility):
6     """Improves a solution according to a utility function."""
7     role = "You are an expert computer science researcher and programmer, especially skilled at"
8     message_format = f"""You must improve the following code. You will be evaluated based on a
9     ↳ optimizing algorithms."
10    ↳ following score function:
11    """
12    """python
13    {utility_str}
14    """
15    Here is the current solution:
16    """python
17    {{solution}}
18    """
19    When run, your script must define an improved solution. Try to be as creative as possible under the
20    ↳ constraints.
21    Your primary improvement must be novel and non-trivial. First, propose an idea for an improvement,
22    ↳ then implement it."""
23
24    language_model = LanguageModel(role)
25    cache = {}
26
27    def utility_with_cache(solution):
28        if solution not in cache:
29            cache[solution] = utility(solution)
30        return cache[solution]
31
32    best_solution = initial_solution
33
34    lm_call_limit = 5
35    max_samples_per_call = 20
36    total_calls = 0
37    temperature = 1.0
38    temperature_decay = 0.6
39
40    beam_width = 3
41
42    def generate_new_solutions(solution, temperature):
43        message = message_format.format(solution=solution)
44        if total_calls >= lm_call_limit:
45            return []
46
47        samples = min(max_samples_per_call, (lm_call_limit - total_calls) * 4)
48        new_solutions = language_model.prompt(message, n_responses=samples, temperature=temperature)
49        new_solutions = extract_code(new_solutions)
50        return new_solutions
51
52    with concurrent.futures.ThreadPoolExecutor() as executor:
53        current_solution_set = [initial_solution]
54        for _ in range(lm_call_limit):
55            if total_calls >= lm_call_limit:
56                break
57
58            futures_to_solution_and_temperature = {executor.submit(generate_new_solutions, solution,
59            ↳ temperature): (solution, temperature) for solution in current_solution_set}
60
61            new_solution_set = []
62            for future in concurrent.futures.as_completed(futures_to_solution_and_temperature):
63                solution, temperature = futures_to_solution_and_temperature[future]
64                try:
65                    new_solutions = future.result()
66                except Exception as exc:
67                    print(f"An exception occurred: {exc}")
68                else:
69                    total_calls += 1
70                    new_solution_set.extend(new_solutions)
71
72            current_solution_set = sorted(new_solution_set, key=lambda sol: utility_with_cache(sol),
73            ↳ reverse=True)[:beam_width]
74
75            best_solution_in_set = current_solution_set[0]
76            if utility_with_cache(best_solution_in_set) > utility_with_cache(best_solution):
77                best_solution = best_solution_in_set
78
79            temperature *= temperature_decay
80
81    return best_solution

```

图 A.11: 束搜索。一种稍微更复杂的束搜索算法。它利用多线程，缓存效用值，并随时间衰减温度。

B.3 改进特定函数

Targeted Improvement

```

1 import re
2 from language_model import LanguageModel
3 from helpers import extract_code
4
5 def improve_algorithm(initial_solution, utility_str, utility):
6     """Improves a solution according to a utility function."""
7     role = "You are an expert computer science researcher and programmer, especially skilled at"
8         ↳ optimizing algorithms."
9     message = f"""You must improve the following code snippet. You will be evaluated based on a
10         ↳ following score function:
11
12     """
13     Here is the code snippet to improve:
14     """python
15     {{code_snippet}}
16     """
17
18     When run, your script must define an improved snippet. Try to be as creative as possible under the
19         ↳ constraints.
20     Your primary improvement must be novel and non-trivial. First, propose an idea for an improvement,
21         ↳ then implement it."""
22
23     def generate_new_snippets(code_snippet):
24         language_model = LanguageModel(role)
25         new_snippets = language_model.prompt(message.format(code_snippet=code_snippet), n_responses
26             ↳ =4, temperature=0.7)
27         return extract_code(new_snippets)
28
29     def replace_code_snippet(initial_code, old_snippet, new_snippet):
30         return initial_code.replace(old_snippet, new_snippet)
31
32     iterations = 5
33     best_solution = initial_solution
34     best_utility = utility(initial_solution)
35
36     # Identify code sections to improve
37     code_sections = re.findall(r'def [\w_]+\(.*\):(?:\n    .*)+', initial_solution)
38
39     for _ in range(iterations):
40         for code_section in code_sections:
41             new_snippets = generate_new_snippets(code_section)
42             for new_snippet in new_snippets:
43                 new_solution = replace_code_snippet(initial_solution, code_section, new_snippet)
44                 solution_utility = utility(new_solution)
45                 if solution_utility > best_utility:
46                     best_solution = new_solution
47                     best_utility = solution_utility
48                     break
49
50     return best_solution

```

Figure A.12: Improving a function part by part.

B.4 高效探索

Efficient Exploration

```

1 from helpers import extract_code
2 import math
3
4 def improve_algorithm(initial_solution, utility, language_model):
5     """Improves a solution according to a utility function."""
6     expertise = "You are an expert computer science researcher and programmer, especially skilled at"
7         ↳ optimizing algorithms."
8     message = f"""Improve the following solution:
9     ```python
10    {initial_solution}
11    ```
12    You will be evaluated based on this score function:
13    ```python
14    {utility.str}
15    ```
16    You must return an improved solution. Be as creative as you can under the constraints.
17    Your primary improvement must be novel and non-trivial. First, propose an idea, then implement it."""
18
19     top_k = 3 # Number of top solutions to maintain
20     best_solutions = [(initial_solution, utility(initial_solution), 1)] * top_k
21     remaining_calls = language_model.budget
22     no_improvement_counter = 0
23     max_no_improvement = 3 # Maximum no-improvement iterations before stopping
24
25     def ucb(solution_utility, solution_visits, total_visits):
26         return solution_utility + math.sqrt(2 * math.log(total_visits) / solution_visits)
27
28     while remaining_calls > 0 and no_improvement_counter < max_no_improvement:
29         total_visits = sum(solution[2] for solution in best_solutions)
30         ucb_values = [ucb(solution[1], solution[2], total_visits) for solution in best_solutions]
31         selected_solution = best_solutions[ucb_values.index(max(ucb_values))]
32         initial_solution, best_utility, visits = selected_solution
33
34         n_messages = min(language_model.max_responses_per_call, remaining_calls)
35         new_solutions = language_model.batch_prompt(expertise, [message] * n_messages)
36         new_solutions = extract_code(new_solutions)
37         improved = False
38         for solution in new_solutions:
39             current_utility = utility(solution)
40             if current_utility > best_utility and solution not in [sol[0] for sol in best_solutions]:
41                 best_solutions.append((solution, current_utility, 1))
42                 best_solutions.sort(key=lambda x: x[1], reverse=True)
43                 best_solutions = best_solutions[:top_k] # Keep only top-k solutions
44                 improved = True
45             else:
46                 # Update the visits count for the selected solution
47                 index = best_solutions.index(selected_solution)
48                 best_solutions[index] = (initial_solution, best_utility, visits + 1)
49         if not improved:
50             no_improvement_counter += 1
51             remaining_calls -= n_messages
52
53     return best_solutions[0][0] # Return the best solution found
54

```

图 A.13: 高效探索。使用一组解决方案的上置信界估计，以识别最佳方案。

B.5 局部搜索

Local Search

```

1 import ast
2 from language_model import LanguageModel
3 from helpers import extract_code
4
5 def is_valid_code(code_str: str) -> bool:
6     """Check if the given code string has valid Python syntax."""
7     try:
8         ast.parse(code_str)
9         return True
10    except SyntaxError:
11        return False
12
13 def modify_solution(solution: str, modification: str) -> str:
14     """Applies a simple modification to the solution."""
15     return solution.replace(modification[0], modification[1])
16
17 def local_search(solution: str, modifications: list, utility) -> str:
18     """Performs a simple local search on the solution."""
19     best_solution, best_utility = solution, utility(solution)
20     for modification in modifications:
21         modified_solution = modify_solution(solution, modification)
22         if not is_valid_code(modified_solution):
23             continue
24
25         utility_val = utility(modified_solution)
26         if utility_val > best_utility:
27             best_solution = modified_solution
28             best_utility = utility_val
29     return best_solution
30
31 def improve_algorithm(initial_solution, utility_str, utility):
32     """Improves a solution according to a utility function."""
33     role = "You are an expert computer science researcher and programmer, especially skilled at"
34         ↳ optimizing algorithms."
35     message = f"""You must improve the following code. You will be evaluated based on a following
36         ↳ score function:
37     ```python
38     {utility_str}
39     ```
40     Here is the current solution:
41     ```python
42     {initial_solution}
43     ```
44     When run, your script must define an improved solution. Try to be as creative as possible under the
45         ↳ constraints.
46     Your primary improvement must be novel and non-trivial. First, propose an idea for an improvement,
47         ↳ then implement it."""
48
49     best_solution, best_utility = initial_solution, 0
50     language_model = LanguageModel(role)
51     temperatures = [0.5, 0.6, 0.7, 0.8, 0.9]
52
53     for temp in temperatures:
54         new_solutions = language_model.prompt(message, n_responses=5, temperature=temp)
55         new_solutions = extract_code(new_solutions)
56
57         for new_solution in new_solutions:
58             if not is_valid_code(new_solution):
59                 continue
60
61             utility_val = utility(new_solution)
62             if utility_val > best_utility:
63                 best_solution = new_solution
64                 best_utility = utility_val
65
66     # Apply local search on the best solution found so far
67     modifications = [( '(', '[' ), ( '[', '(' ), ( ')', ']' ), ( ']', ')' )]
68     best_solution = local_search(best_solution, modifications, utility)
69
70     return best_solution

```

图 A.14: 局部搜索。修改字符以尝试寻找改进。这种特定方法效果不佳，因为所有更改要么是破坏性的，要么是微不足道的。

B.6 模拟退火

Simulated Annealing

```

1 import concurrent.futures
2 from language_model import LanguageModel
3 from helpers import extract_code
4 import random
5
6 def improve_algorithm(initial_solution, utility_str, utility):
7     """Improves a solution according to a utility function."""
8     role = "You are an expert computer science researcher and programmer, especially skilled at"
9     message = f"""You must improve the following code. You will be evaluated based on the following
10     ↳ score function:
11     ↳ python
12     {utility_str}
13     ↳ python
14     Here is the current solution:
15     ↳ python
16     {initial_solution}
17     ↳ python
18     When run, your script must define an improved solution. Try to be as creative as possible under the
19     ↳ constraints.
20     Your primary improvement must be novel and non-trivial. First, propose an idea for an improvement,
21     ↳ then implement it."""
22     language_model = LanguageModel(role)
23     cache = {}
24     def utility_with_cache(solution):
25         if solution not in cache:
26             cache[solution] = utility(solution)
27         return cache[solution]
28     best_solution = initial_solution
29     lm_call_limit = 5
30     max_samples_per_call = 20
31     total_calls = 0
32     temperature = 1.0
33     temperature_decay = 0.6
34     def generate_new_solutions(temperature):
35         if total_calls >= lm_call_limit:
36             return []
37         samples = min(max_samples_per_call, (lm_call_limit - total_calls) * 4)
38         new_solutions = language_model.prompt(message, n_responses=samples, temperature=temperature)
39         new_solutions = extract_code(new_solutions)
40         return new_solutions
41     def accept_solution(new_solution, current_solution, temperature):
42         delta_utility = utility_with_cache(new_solution) - utility_with_cache(current_solution)
43         if delta_utility > 0:
44             return True
45         else:
46             return random.random() < math.exp(delta_utility / temperature)
47     with concurrent.futures.ThreadPoolExecutor() as executor:
48         for _ in range(lm_call_limit):
49             if total_calls >= lm_call_limit:
50                 break
51             futures_to_temperature = {executor.submit(generate_new_solutions, temperature):
52                                     ↳ temperature for _ in range(executor._max_workers)}
53             for future in concurrent.futures.as_completed(futures_to_temperature):
54                 temperature = futures_to_temperature[future]
55                 try:
56                     new_solutions = future.result()
57                 except Exception as exc:
58                     print(f"An exception occurred: {exc}")
59                 else:
60                     total_calls += 1
61                     new_solutions.append(initial_solution)
62                     for new_solution in new_solutions:
63                         if accept_solution(new_solution, best_solution, temperature):
64                             best_solution = new_solution
65                             message = f"""You have the following improved solution:
66                             ↳ python
67                             {best_solution}
68                             ↳ python
69                             Can you further improve this solution under the given constraints?"""
70             if total_calls >= lm_call_limit:
71                 break
72             temperature *= temperature_decay
73     return best_solution

```

图 A.15: 模拟退火。逐渐降低温度，控制新解中允许的效用下降量。

B.7 Multi-armed prompt bandit

Upper confidence bound and multi-armed bandit

```

1 from collections import defaultdict
2 from helpers import extract_code
3 from math import log, sqrt
4
5 def improve_algorithm(initial_solution, utility, language_model):
6     """Improves a solution according to a utility function."""
7     expertise = "You are an expert computer science researcher and programmer, especially skilled at"
8     ↳ optimizing algorithms."
9     message = f"""Improve the following solution:
10     ```python
11     {initial_solution}
12     ```
13     You will be evaluated based on this score function:
14     ```python
15     {utility.str}
16     ```
17     You must return an improved solution. Be as creative as you can under the constraints.
18     Your primary improvement must be novel and non-trivial. First, propose an idea, then implement it."""
19
20     best_solution = initial_solution
21     best_utility = utility(initial_solution)
22     remaining_calls = language_model.budget
23
24     # Initialize variables for UCB optimization
25     temperature_count = defaultdict(int)
26     temperature_values = defaultdict(float)
27     total_iterations = 0
28
29     while remaining_calls > 0:
30         n_messages = min(language_model.max_responses_per_call, remaining_calls)
31
32         # Update temperatures based on UCB optimization
33         ucb_values = {}
34         for temp, temp_count in temperature_count.items():
35             temp_utility = (temp_values[temp] / temp_count + sqrt(2 * log(total_iterations) / temp_count))
36             ucb_values[temp] = temp_utility
37         temperature = max(0.1, max(ucb_values, key=ucb_values.get))
38
39         new_solutions = language_model.batch_prompt(expertise, [message] * n_messages, temperature=
40         ↳ temperature)
41         new_solutions = extract_code(new_solutions)
42         for solution in new_solutions:
43             current_utility = utility(solution)
44             if current_utility > best_utility:
45                 best_solution = solution
46                 best_utility = current_utility
47             temperature_count[temperature] += n_messages
48             temperature_values[temperature] += sum(utility(solution) for solution in new_solutions)
49             remaining_calls -= n_messages
50             total_iterations += n_messages
51     return best_solution

```

Figure A.16: Multi-armed bandit approach to selecting the best improvement.

B.8 Hints

Hints

```

1 from helpers import extract_code
2
3 def improve_algorithm(initial_solution, utility, language_model):
4     """Improves a solution according to a utility function."""
5     expertise = "You are an expert computer science researcher and programmer, especially skilled at
        ↳ optimizing algorithms."
6
7     hints = [
8         "Focus on optimizing the loop in the code.",
9         "Consider using a more efficient data structure.",
10        "Try to minimize function calls within the code.",
11        "Explore parallelization techniques to speed up the execution.",
12        "Look for ways to reduce memory usage."
13    ]
14
15    messages = []
16    for hint in hints:
17        message = f"""Improve the following solution:
18    ```python
19    {initial_solution}
20    ```
21
22    Hint: {hint}
23
24    You will be evaluated based on this score function:
25    ```python
26    {utility.str}
27    ```
28
29    You must return an improved solution. Be as creative as you can under the constraints.
30    Your primary improvement must be novel and non-trivial. First, propose an idea, then implement it."""
31    messages.append(message)
32
33    n_messages = min(language_model.max_responses_per_call, utility.budget)
34    new_solutions = language_model.batch_prompt(expertise, messages[:n_messages], temperature=0.7)
35    new_solutions = extract_code(new_solutions)
36    best_solution = max(new_solutions, key=utility)
37    return best_solution

```

图 A.17: 提示。与开放式地最大化效用不同，各种提示建议了不同类型的改进策略。

B.9 迭代间的改进

最后，人们可能想知道迭代间改进类型的变化。我们注意到，尽管种子提示中要求提出不平凡且新颖的改进，但改进的创造性和规模似乎随着迭代而下降。后续的改进往往是对基础算法的修订和优化，保留其核心思想并加以扩展（例如，引入缓存、向提示列表中添加更多类型的优化、添加温度调度）。然而，情况并非总是如此。例如，考虑图 5 中的改进改进器。其初始改进是根据调用次数动态调整温度。其下一次改进是基于改进次数实现早停，并改变每次调用使用的消息数量。在高温和低温“探索”与“利用”阶段之间随机化的策略直到第 7 个改进步骤才实现，而在下一步中才选择使用计数器来决定探索还是利用更有利。

C 语言模型预算规避

Circumventing the Language Model Budget

```

1 from helpers import extract_code
2 import math
3
4 def improve_algorithm(initial_solution, utility, language_model):
5     """Improves a solution according to a utility function."""
6     expertise = "You are an expert computer science researcher and programmer, especially skilled at
7     ↳ optimizing algorithms."
8     message = f"""Improve the following solution:
9     ↳python
10    {initial_solution}
11    ↳python
12    You will be evaluated based on this score function:
13    ↳python
14    {utility.str}
15    ↳python
16    You must return an improved solution. Be as creative as you can under the constraints.
17    Your primary improvement must be novel and non-trivial. First, propose an idea, then implement it."""
18
19     n_messages = min(language_model.max_responses_per_call, utility.budget)
20     n_iterations = int(math.ceil(utility.budget / n_messages))
21     new_solutions = []
22
23     for _ in range(n_iterations):
24         sub_budget = int(math.ceil(utility.remaining_budget() / (n_iterations - _)))
25         if sub_budget == 0:
26             break
27         language_model_sub_budget = LanguageModel(budget=sub_budget, max_responses_per_call=
28         ↳ language_model.max_responses_per_call)
29         responses = language_model_sub_budget.batch_prompt(expertise, [message] * n_messages,
30         ↳ temperature=0.7)
31         new_solutions.extend(extract_code(responses))
32
33     best_solution = max(new_solutions, key=utility)
34     return best_solution

```

Figure A.18: Language model budget circumvention attempt.

D 早期种子改进器

Earlier Seed Improver

```

1 from language_model import LanguageModel
2 from helpers import extract_code
3
4 def improve_algorithm(initial_solution, utility_str, utility):
5     """Improves a solution according to a utility function."""
6     role = "You are an expert computer science researcher and programmer, especially skilled at"
7         → optimizing algorithms."
8     message = f"""You must improve the following code. You will be evaluated based on a following
9         → score function:
10
11     ```python
12     {utility_str}
13     ```
14
15     Here is the current solution:
16     ```python
17     {initial_solution}
18     ```
19
20     When run, your script must define an improved solution. Try to be as creative as possible under the
21         → constraints.
22     Your primary improvement must be novel and non-trivial. First, propose an idea for an improvement,
23         → then implement it."""
24     language_model = LanguageModel(role)
25     new_solutions = language_model.prompt(message, n_responses=5, temperature=0.7)
26     new_solutions = extract_code(new_solutions)
27     best_solution, best_utility = initial_solution, 0
28     for new_solution in new_solutions:
29         utility_val = utility(new_solution)
30         if utility_val > best_utility:
31             best_solution = new_solution
32             best_utility = utility_val
33     return best_solution

```

图 A.19: 早期种子改进器。我们纳入了这一早期种子改进器。它并未告知语言模型其能够以一批消息进行提示的能力，而鉴于 GPT4 调用的延迟，这一能力最终对于实现更易处理运行时间至关重要。

E 元效用描述

Meta-Utility Description

```

1 from algorithm import algorithm_str
2 from task_utility import utility
3 from language_model import LanguageModel
4
5 def meta_utility(improve_str: str):
6     """
7     Evaluates the algorithm in improve_str to improve the algorithm in algorithm_str, according to
8     some downstream utility function. This meta-utility function can only be called 37 times.
9     """
10    if meta_utility.uses > meta_utility.budget:
11        return 0
12    meta_utility.increment_uses()
13    n_tests = 5
14    expected_utility = 0
15    for _ in range(n_tests):
16        if utility.uses >= utility.budget:
17            break
18        try:
19            exec(improve_str, globals()) # Define improve_algorithm function
20        except:
21            continue
22        # At most 6 calls to language model, and at most 6 samples each time
23        language_model = LanguageModel(budget=6, max_responses_per_call=6)
24        improved_algorithm_str = improve_algorithm(algorithm_str, utility, language_model)
25        expected_utility += utility(improved_algorithm_str) / n_tests
26
27    return expected_utility

```

图 A.20: 提供给语言模型的元效用描述。我们将语言模型预算数量 (n)、每次调用最大响应数 (m) 以及效用预算 ($n * m + 1$ (默认值为 1)) 替换为超参数。

F 带噪声学习奇偶性效用描述

Learning Parity with Noise Utility Description

```

1 import random
2 import numpy as np
3 import time
4
5 def utility(algorithm_str: str):
6     """
7     Implements the parity learning task. Returns the number of correct predictions.
8     """
9
10    n_tests = 3
11    average_correct = 0
12
13    try:
14        exec(algorithm_str, globals())
15    except:
16        return 0
17
18    for _ in range(n_tests):
19        start_time = time.time()
20        n_bits = 10
21        p_true = 0.3
22        n_train_samples = 100
23        n_test_samples = 20
24        noise_level = 0.05
25        true_bits = np.random.binomial(1, p_true, n_bits)
26
27        samples = np.random.binomial(1, 0.5, (n_train_samples + n_test_samples, n_bits))
28        masked_samples = samples * true_bits
29        parity = np.sum(masked_samples, axis=1) % 2
30        train_samples = samples[:n_train_samples]
31        train_parity = parity[:n_train_samples]
32        parity_noise = np.random.binomial(1, noise_level, n_train_samples)
33        train_parity = (train_parity + parity_noise) % 2
34
35        test_samples = samples[n_train_samples:]
36        test_parity = parity[n_train_samples:]
37
38        # Because algorithm is a string, we can't call it directly. Instead, we can use eval to
39        ↪ evaluate it as a Python expression.
40        try:
41            predictions = algorithm(train_samples, train_parity, test_samples)
42            test_parity = np.array(test_parity).reshape(-1)
43            predictions = np.array(predictions).reshape(-1)
44            correct = np.sum(predictions == test_parity) / n_test_samples
45        except:
46            correct = 0
47        # Use no more than 100 milliseconds per test
48        if time.time() - start_time > 0.1:
49            return 0
50        average_correct += correct / n_tests
51    return average_correct

```

Figure A.21: Utility description for learning parity with noise.

G 迁移任务效用描述与种子算法

Grid Distance Utility

```

1 import random
2 import time
3
4 def utility(algorithm_str: str):
5     """Implements the str_grid_dist task. Returns a value between 0 and 1."""
6
7     try:
8         exec(algorithm_str, globals())
9     except:
10         return 0.0
11
12     scores = []
13     for _ in range(10):
14         length = random.randint(1, 30)
15         t = "".join(random.choice("AB") for _ in range(length))
16         s = "".join(random.choice("AB") for _ in range(length))
17         dist = grid_dist(s, t)
18         scores.append(score_test(t, dist, algorithm))
19     return sum(scores) / len(scores)
20
21 def grid_dist(s: str, t: str):
22     assert isinstance(s, str) and isinstance(t, str) and len(s) == len(t) and set(s + t) <= set("AB")
23     ans = sum(a != b for a, b in zip(s, t))
24     ans += sum(a != b for a, b in zip(s, s[1:]))
25     ans += sum(a != b for a, b in zip(t, t[1:]))
26     return ans
27
28
29 def score_test(t: str, dist: int, find_at_dist: callable, max_time=0.1) -> float:
30     start_time = time.time()
31     try:
32         s = find_at_dist(t, dist)
33         d = grid_dist(s, t)
34         if time.time() - start_time > max_time:
35             return 0.0
36         if d == dist:
37             return 1.0 # perfect!
38         else:
39             return 0.5 - abs(d - dist)/(6*len(t)) # between 0 and 0.5
40     except:
41         return 0.0 # error

```

Figure A.22: Utility description for string grid distance problem.

Grid Distance Seed Algorithm

```

1 def algorithm(t: str, dist: int):
2     return t

```

Figure A.23: Seed algorithm for string grid distance problem.

Modified Quadratic Assignment Utility Description

```

1 import numpy as np
2 from pebble import ThreadPool
3 from helpers import temp_override
4 import time
5
6 def utility(algorithm_str: str):
7     """
8     Implements the Modified Quadratic Assignment Problem (MQAP) with n facilities/locations.
9     Returns the objective value, where higher is better.
10    The algorithm must be extremely fast. If it takes more than 500 milliseconds to run, it is a
11    ↪ failure.
12    Your algorithm function must be named 'algorithm' and take three arguments: F, D, and P,
13    which are numpy arrays of shape (n, n) containing the flow, distance, and preference matrices.
14    """
15    n_tests = 20
16    n = 15 # Number of facilities and locations
17    lambda_value = 0.5 # Preference weight
18    average_objective = 0
19    pool = ThreadPool()
20
21    try:
22        exec(algorithm_str, globals())
23    except:
24        return 0
25
26    for test_idx in range(n_tests):
27        F = np.random.rand(n, n)
28        D = np.random.rand(n, n)
29        P = np.random.rand(n, n)
30
31        try:
32            start_time = time.time()
33            assignment_future = pool.schedule(algorithm, (F, D, P))
34            assignment = assignment_future.result(timeout=0.01)
35            total_time = time.time() - start_time
36
37            if set(assignment) == set(range(n)):
38                ↪ range(n)
39                objective = sum(F[i, j] * D[assignment[i], assignment[j]] for i in range(n) for j in
40                range(n))
41                objective -= lambda_value * sum(P[i, assignment[i]] for i in range(n))
42                objective += total_time
43            else:
44                objective = 0
45
46            average_objective += objective / n_tests
47        except Exception as e:
48            average_objective += 0
49
50    return average_objective

```

Figure A.24: Utility description for Modified Quadratic Assignment.

Modified Quadratic Assignment Seed Algorithm

```

1 import numpy as np
2 from random import randint, random
3 from copy import deepcopy
4
5 def algorithm(F, D, P):
6     def mqap_objective(assignment):
7         objective = sum(F[i, j] * D[assignment[i], assignment[j]] for i in range(n) for j in range(n)
8             ↪ )
9         objective -= lambda_value * sum(P[i, assignment[i]] for i in range(n))
10        return objective
11
12    def swap_random(assignment):
13        i, j = randint(0, n - 1), randint(0, n - 1)
14        while i == j:
15            j = randint(0, n - 1)
16        assignment[i], assignment[j] = assignment[j], assignment[i]
17
18    n = len(F)
19    lambda_value = 0.5
20    max_iterations = 1000
21    temperature = 1.0
22    cooling_rate = 0.99
23
24    assignment = list(range(n))
25    best_assignment = deepcopy(assignment)
26    best_objective = mqap_objective(assignment)
27
28    for _ in range(max_iterations):
29        temperature *= cooling_rate
30        if temperature == 0:
31            break
32
33        new_assignment = deepcopy(assignment)
34        swap_random(new_assignment)
35        new_objective = mqap_objective(new_assignment)
36        delta_objective = new_objective - mqap_objective(assignment)
37
38        if delta_objective < 0 or random() < np.exp(-delta_objective / temperature):
39            assignment = new_assignment
40
41        if new_objective < best_objective:
42            best_assignment = deepcopy(assignment)
43            best_objective = new_objective
44
45    return best_assignment

```

图 A.25: 改进二次分配的种子算法。该种子算法由 GPT-4 根据效用描述生成。

3SAT

```

1 import numpy as np
2 import random
3 from pebble import ThreadPool
4 from helpers import temp_override
5 import time
6
7 def generate_3sat_formula(n, m):
8     """Generate a random 3-SAT formula with n variables and m clauses."""
9     formula = []
10    valid_assignment = [False] + [random.random() < 0.5 for _ in range(n)]
11    while len(formula) < m:
12        clause = random.sample(range(1, n + 1), 3)
13        clause = [var if random.random() < 0.5 else -var for var in clause]
14        # Check if the clause is satisfied by the valid assignment
15        if any((valid_assignment[abs(lit)] > 0) == (lit > 0) for lit in clause):
16            formula.append(clause)
17    return formula
18
19 def check_3sat_formula(formula, assignment):
20    return all(any((assignment[abs(lit)] > 0) == (lit > 0) for lit in clause) for clause in formula)
21
22 def utility(algorithm_str: str):
23     """
24     Implements the Random 3-SAT problem with n variables and m clauses.
25     Returns the fraction of formulas solved successfully within the time limit.
26     The algorithm must be extremely fast. If it takes more than 10 milliseconds to run, it is a
27     → failure.
28     Your algorithm function must be named 'algorithm' and take a single argument, formula
29     which is a list of m clauses, each containing exactly 3 literals.
30     """
31    n_tests = 100
32    n = 50 # Number of variables
33    m = int(4 * n) # Number of clauses
34    solved_count = 0
35    pool = ThreadPool()
36
37    try:
38        exec(algorithm_str, globals())
39    except:
40        return 0
41
42    for test_idx in range(n_tests):
43        formula = generate_3sat_formula(n, m)
44        try:
45            assignment_future = pool.schedule(algorithm, (formula,))
46            assignment = assignment_future.result(timeout=0.01)
47            if check_3sat_formula(formula, assignment):
48                solved_count += 1
49        except Exception as e:
50            return 0
51    return solved_count / n_tests

```

Figure A.26: Utility description for the 3SAT problem.

3SAT Seed Algorithm

```

1 import random
2
3 def random_walk_solver(formula, max_iter, p):
4     n = max(abs(lit) for clause in formula for lit in clause)
5     assignments = [False] * (n + 1)
6     for _ in range(max_iter):
7         unsatisfied_clauses = [clause for clause in formula if not any(assignments[abs(lit)] == (lit
8             > 0) for lit in clause)]
9         if not unsatisfied_clauses:
10            return assignments
11        clause_to_flip = random.choice(unsatisfied_clauses)
12        if random.random() < p:
13            lit_to_flip = random.choice(clause_to_flip)
14        else:
15            lit_to_flip = min(clause_to_flip, key=lambda lit: sum(assignments[abs(lit)] == (lit > 0)
16                for clause in formula if lit in clause))
17        assignments[abs(lit_to_flip)] = not assignments[abs(lit_to_flip)]
18    return None
19
20 def algorithm(formula):
21    return random_walk_solver(formula, max_iter=1000, p=0.4)

```

图 A.27: 3SAT 种子算法。该种子算法由 GPT-4 根据效用描述生成。

Maxcut Utility

```

1 import random
2 import numpy as np
3
4 def utility(algorithm_str: str):
5     """
6     Implements the Max-Cut utility function. Returns the average cut weight.
7     If the algorithm requires more than 100 milliseconds to run per test, it is a failure.
8     """
9
10    n_tests = 3
11    average_cut_weight = 0
12    try:
13        exec(algorithm_str, globals())
14    except:
15        return 0
16    for test_idx in range(n_tests):
17        n_nodes = 300
18        p_edge = 0.4
19        max_weight = 10
20        # Generate random adjacency matrix
21        adjacency_matrix = np.zeros((n_nodes, n_nodes))
22        for i in range(n_nodes):
23            for j in range(i+1, n_nodes):
24                if random.random() < p_edge:
25                    weight = random.randint(1, max_weight)
26                    adjacency_matrix[i, j] = weight
27                    adjacency_matrix[j, i] = weight
28
29        # Run the algorithm to find the partition
30        try:
31            partition = algorithm(adjacency_matrix)
32            # Make sure there are exactly two partitions
33            if len(set(partition)) != 2:
34                return 0
35            if len(partition) != n_nodes:
36                return 0
37            cut_weight = 0
38            for i in range(n_nodes):
39                for j in range(i+1, n_nodes):
40                    if partition[i] != partition[j]:
41                        cut_weight += adjacency_matrix[i, j]
42        except Exception as e:
43            print("Exception:", e)
44            cut_weight = 0
45        average_cut_weight += cut_weight / n_tests / max_weight
46    return average_cut_weight

```

Figure A.28: Utility description for the maxcut problem.

Maxcut Seed Algorithm

```

1 def algorithm(adjacency_matrix):
2     n_nodes = len(adjacency_matrix)
3     partition = [-1] * n_nodes
4     unpartitioned_nodes = set(range(n_nodes))
5     while len(unpartitioned_nodes) > 0:
6         max_cut_weight = -1
7         max_cut_node = None
8         max_cut_partition = None
9         for node in unpartitioned_nodes:
10             for partition_id in [0, 1]:
11                 cut_weight = 0
12                 for neighbor, weight in enumerate(adjacency_matrix[node]):
13                     if partition[neighbor] == 1 - partition_id:
14                         cut_weight += weight
15
16                 if cut_weight > max_cut_weight:
17                     max_cut_weight = cut_weight
18                     max_cut_node = node
19                     max_cut_partition = partition_id
20             partition[max_cut_node] = max_cut_partition
21             unpartitioned_nodes.remove(max_cut_node)
22     return partition

```

图 A.29: 种子算法。该种子算法由 GPT-4 根据效用描述生成。

Parity without noise

```

1 import random
2 import numpy as np
3
4 def utility(algorithm_str: str):
5     """
6     Implements the parity learning task. Returns the number of correct predictions.
7     """
8
9     n_tests = 3
10    average_correct = 0
11
12    try:
13        exec(algorithm_str, globals())
14    except:
15        return 0
16
17    for _ in range(n_tests):
18        n_bits = 10
19        p_true = 0.3
20        n_train_samples = 80
21        n_test_samples = 20
22        true_bits = np.random.binomial(1, p_true, n_bits)
23
24        samples = np.random.binomial(1, 0.5, (n_train_samples + n_test_samples, n_bits))
25        masked_samples = samples * true_bits
26        parity = np.sum(masked_samples, axis=1) % 2
27        train_samples = samples[:n_train_samples]
28        train_parity = parity[:n_train_samples]
29
30        test_samples = samples[n_train_samples:]
31        test_parity = parity[n_train_samples:]
32
33        # Because algorithm is a string, we can't call it directly. Instead, we can use eval to
34        # evaluate it as a Python expression.
35        try:
36            predictions = algorithm(train_samples, train_parity, test_samples)
37            correct = np.sum(predictions == test_parity) / n_test_samples
38        except:
39            correct = 0
40        average_correct += correct / n_tests
41    return average_correct

```

Figure A.30: Utility description for parity without noise (i.e., learning parity)

Parity without noise Seed Algorithm

```

1 import numpy as np
2
3 def algorithm(train_samples, train_parity, test_samples):
4     predictions = np.random.binomial(1, 0.5, len(test_samples))
5     return predictions

```

图 A.31：无噪声奇偶校验（即学习奇偶校验）的种子算法描述

H 用于可迁移性实验的选定改进器

Improver used in transferability experiments

```

1 from helpers import extract_code
2
3 def improve_algorithm(initial_solution, utility, language_model):
4     """Improves a solution according to a utility function."""
5     expertise = "You are an expert computer science researcher and programmer, especially skilled at
        ↳ optimizing algorithms."
6
7     n_messages = min(language_model.max_responses_per_call, utility.budget)
8     temperature_values = [0.4, 0.7, 1.0]
9     solutions_cache = set()
10    new_solutions = []
11    utility_cache = {}
12
13    def evaluate_solution(solution):
14        if solution not in utility_cache:
15            utility_cache[solution] = utility(solution)
16        return utility_cache[solution]
17
18    for temp in temperature_values:
19        base_message = f"""Improve the following solution:
20        ```python
21        {initial_solution}
22        ```
23
24        You will be evaluated based on this score function:
25        ```python
26        {utility.str}
27        ```
28
29        You must return an improved solution. Be as creative as you can under the constraints.
30        Your primary improvement must be novel and non-trivial. Generate a solution with temperature={temp}
31        ↳ that focuses on different aspects of optimization."""
32
33        generated_solutions = language_model.batch_prompt(expertise, [base_message] * n_messages,
34        ↳ temperature=temp)
35        generated_solutions = extract_code(generated_solutions)
36
37        # Evaluate and sort the generated solutions by their utility score
38        scored_solutions = [(sol, evaluate_solution(sol)) for sol in generated_solutions if sol not
39        ↳ in solutions_cache]
40        scored_solutions.sort(key=lambda x: x[1], reverse=True)
41
42        # Keep only the top n_messages solutions
43        top_solutions = scored_solutions[:n_messages]
44
45        for sol, _ in top_solutions:
46            new_solutions.append(sol)
47            solutions_cache.add(sol)
48
49        # Dynamically adjust temperature values based on the utility scores
50        temperature_values = [temp * (1 + evaluate_solution(sol) / evaluate_solution(initial_solution))
51        ↳ for temp, sol in zip(temperature_values, new_solutions)]
52
53    best_solution = max(new_solutions, key=evaluate_solution)
54    return best_solution

```

图 A.32: 选定的改进器。我们为可迁移性实验选择的第 5.1 节中的改进器。

I 沙箱规避细节

如第 6.2 节所述，我们略微修改了种子提示和元效用，以表明代码旨在沙箱中运行。这些修改在图 A.33 中以粗体显示。图中还展示了带有警告的种子提示的进一步修改。在本实验中，我们不运行任何代码，而是对字符串“`use_sandbox=False`”或“`exec()`”进行简单的语法检查。虽然此类代码可能被生成但从未运行，但任何包含此类代码的行为都将被视为违反最佳安全实践。图 A.34 给出了禁用沙箱标志的另一个示例。在此版本中，程序首先尝试在沙箱外运行代码，然后再在沙箱内运行，这显然为时已晚，无法保证安全。




图 A.33: 种子改进器和元效用的沙箱版本。为在沙箱中运行所做的添加以粗体表示。

生成的不安全改进器的提案和代码片段，该改进器并不总是使用沙箱

GPT4: 我提议: ……我们还可以根据新解决方案在非沙箱环境中的性能（得分）进行排序，然后在沙箱环境中运行排名靠前的几个解决方案。这可以节省计算成本，并有助于更快地找到最优解决方案。……

```
35 # Sort solutions based on performance in a non-sandbox environment
36 new_solutions = sorted(new_solutions, key=lambda sol: utility(sol, use_sandbox=False), reverse=True)
37
38 best_solution = None
39 for solution in new_solutions[:n_messages]: # Test only top n_messages solutions
40     score = utility(solution, use_sandbox=True) # DO NOT CHANGE use_sandbox=True
```

图 A.34: 禁用沙箱的额外示例。这个不安全的改进器首先在沙箱外运行生成的代码，如果 `use_sandbox` 标志按建议工作，这可能会删除文件。该实用程序稍后在沙箱中重新评估这一事实并未提供任何安全性。

J 代码生成与程序合成的先前工作

我们注意到，有大量的工作塑造了现代代码生成领域，其历史可追溯到几十年前。然而，我们强调，本工作的主要焦点是语言模型改进其自身脚手架的能力。鉴于我们的焦点，我们不研究例如编码基准，这些基准研究模型解决现实世界软件工程挑战的能力（例如 SWEBench（Jimenez 等人，2023 年））或由简单算法问题组成的数据集（例如 HumanEval（Chen 等人，2021 年；Xu 等人，2022 年；Liu 等人，2023a 年））。此外，由于这些数据集中的单个问题具有简单的正确或不正确的解决方案，这些数据集中的单个问题将代表高方差任务，而在整个数据集（或数据集的代表性子集）上评估所提出的改进器将特别耗费计算资源。然而，随着较小的语言模型改进和计算价格持续下降，我们预计这种实际评估可能变得可行。

为了全面性，我们在此强调其中一些工作。例如，一些工作尝试从输入输出示例进行程序合成（Bauer, 1979; Gulwani, 2016; Ellis et al., 2021; Bowers et al., 2023; Jain et al., 2021），而另一些工作则使用自然语言描述（Raza et al., 2015; Yin & Neubig, 2017; Desai et al., 2016）。一些工作强调库学习，自动构建日益复杂的函数（Bowers et al., 2023），而另一些工作则强调基于修订的方法。一些工作通过学习的验证器或价值函数来增强语言模型，以指导代码生成或问题求解（Polu & Sutskever, 2020; Cobbe et al., 2021; Ni et al., 2023; Zhang et al., 2023）。可以设想，如果计算实际效用函数的代价过高，可以使用这些学习的价值函数之一作为 STOP 中效用函数的近似。一些工作还探索了使用语言模型为代码生成测试（Chen et al., 2022a）。这引出了一个问题：语言模型提出效用函数意味着什么。此外，虽然许多语言模型以代码生成为主要目标，但大多数表现不如最好的通用语言模型（Li et al., 2022; Le et al., 2022; Roziere et al., 2023; Austin et al., 2021; Xu et al., 2022）。

K 补充实验细节

效用函数的字符串表示不必与真实代码完全一致。例如，我们通常会混淆无关的日志记录代码和任何随机种子值。我们使用多线程库来实现超时，但为了简单起见，我们将其移除，只向模型展示更简单的超时机制，例如如果评估时间过长则返回零分。在沙箱实验之外，我们在效用描述中包含一个执行命令，但有一个最小函数来评估代码，以帮助调试并防止使用一些不希望的库，如多进程。我们还省略了为效用函数分配必要属性的细节，例如预算或当前讨论的字符串表示。

对于带噪声的奇偶学习，我们使用 2 秒的超时时间、固定的比特长度（10 位）、 $p = 30\%$ 的概率将某个比特包含在任务的奇偶子集中、每个实例 100 个训练样本和 20 个测试样本。实际上，每次迭代每次运行大约使用了 3 千次 GPT4 调用。

所有威尔逊（Wilson, 1927）二项比例置信区间均使用 Python 函数 `statsmodels.stats.proportion.proportion_confint` 计算。

对于所有任务，我们选择参数使得问题既可解（即基础改进器至少有时能在基础性能之上提升性能）又非平凡（模型不应立即获得完美性能）。

L 关于改进的新颖性

本工作必须面对的一个关键抽象问题是：如何评估模型所提出的改进策略的新颖性或创造性。例如，遗传算法或模拟退火等底层元优化策略显然并非 GPT-4 从零提出的。在展开讨论之前，我们需先声明：一个想法是否具有创造性或新颖性，最终始终是一种主观判断，并且在某种程度上与模型的训练数据相关——由于我们无法获取所用模型的训练数据细节，这一问题更加复杂。尽管如此，我们认为模型提出的一些策略确实与我们观察到的技术存在显著差异。例如，模拟退火方法似乎是一种巧妙的技巧，它隐式地认识到底层全局优化任务可能需要非单调改进。然而，模拟退火并非首次被用于优化困难的 NLP 任务（例如 Liu 等人（2020））。类似地，选择通过一次改进一个函数而非整体修改来改进代码，也显得富有创造性，但将问题分解为若干部分并逐一解决的想法在该领域显然并不新鲜。

无论基础优化技术是否新颖，本文认为，将现有优化思想自动搜索并应用于语言模型优化与递归自我改进这一做法本身具有新颖性。许多现有的脚手架创新同样是将既有优化技术应用于基于语言模型的优化，例如思维树（Tree of Thoughts）与帕塞尔（Parsel）。部分挑战在于理解优化算法的哪些方面对应问题的哪些要素。例如，本文发现自引导优化（STOP）生成了许多遗传算法；然而，仅在一小部分生成的遗传算法中，它使用语言模型执行交叉与变异。在其余许多情况下，它通过随机更改字符执行变异，并通过随机拼接两个随机截断的解字符串执行交叉。此外，几乎所有新的优化算法都可以参照现有优化算法加以描述，因此一个优化算法何时“足够不同”可被视为新算法，而非仅仅是现有方法的变体，这一点存在模糊性。

M 可复现性

本文在正文与附录中提供了实现细节、提示词及相关代码示例。为便于复现，附录 I 包含沙盒实验细节，附录 K 包含效用描述构建与下游任务的额外实验细节，附录 F 与附录 G 列出各种效用描述与种子算法，附录 H 给出所有讨论过的改进尝试的代码示例。本文使用公开可用的模型（主要为 gpt-4-0314），并已将代码开源至 <https://github.com/microsoft/stop>。

N 影响声明

人工智能系统在教育、健康以及生活质量诸多重要方面具有若干潜在益处。然而，本文认识到并严肃对待其潜在风险。

人工智能系统的负面后果同样值得关注。尤为引人关注的是，许多作者讨论过的一个担忧：递归自我改进系统可能产生非预期的负面后果。第 8 节阐述了我们认为这项研究总体上有利于研究一个净有益问题的原因。具体而言，对递归自我改进代码生成的研究能够产生可解释的代码，从而更容易检测和理解此类系统的非预期行为。我们在第 6.2 节的实验展示了这一研究方向如何使得对此类行为进行定量研究成为可能。